

BIG DATA with HADOOP

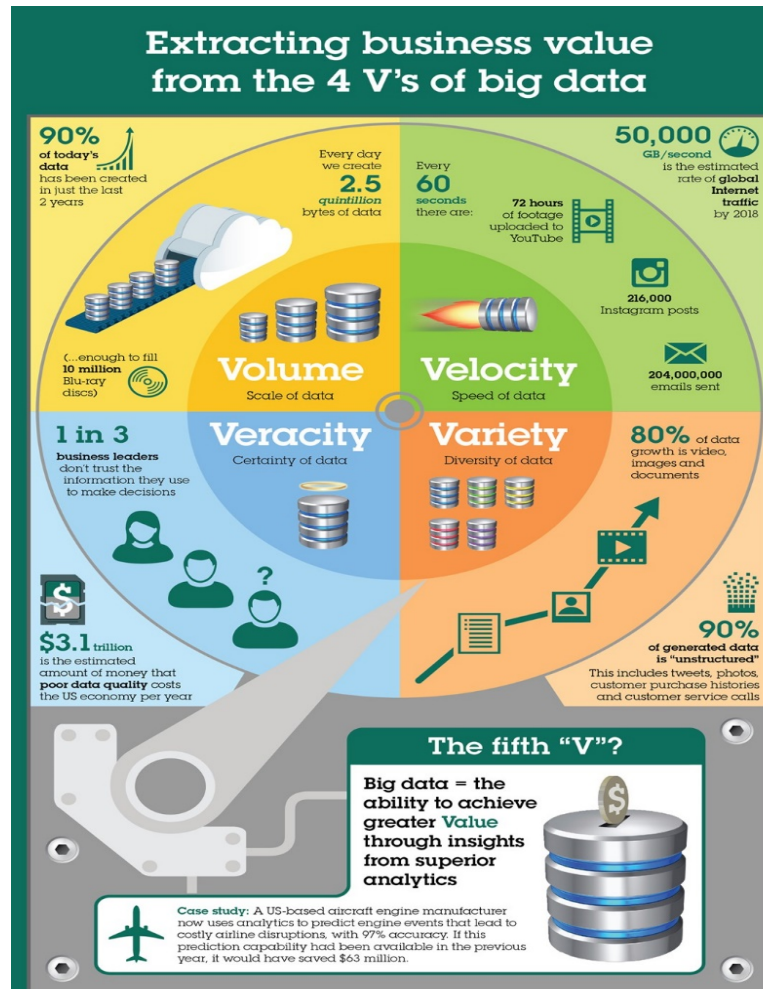


What is big data?

- “Every day, we create 2.5 quintillion bytes of data — so much that 90% of the data in the world today has been created in the last two years alone. This data comes from everywhere: sensors used to gather climate information, posts to social media sites, digital pictures and videos, purchase transaction records, and cell phone GPS signals to name a few.

This data is “big data.”

The 5V's revisited



Unlock the value of your big data.
Start here: ibm.co/technologyplatform

IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the web at "Copyright and trademark information" at www.ibm.com/legal/copytrade.shtml.

IBM

Big Data Database Types

- Key Value Stores
 - Riak(Open Source based on DynamoDB) , Redis, MemcacheD DB, Amazon DynamoDB(not open source), Project Voldemort (Open Source based on DynamoDB)
- Document Databases
 - MongoDB, CouchDB, Terrastore, OrientDB, RavenDB, Lotus Notes storage Facility
- Column Family Databases
 - Cassandra, HBase(Open Source based on Big Table) , Hypertable, Google Big Table(not open source),
- Graph Databases
 - neo4j, Infinite Graph, OrientDB, FlockDB
- **Hadoop Distributed File System (HDFS) – Hadoop**
 - **Does not fit the NoSQL categories as it is a file system but important in this Big Data World**

Apache Hadoop Project

- The Apache Hadoop software library is a framework that allows for the distributed processing of large data sets across clusters of computers using a simple programming model
- <http://hadoop.apache.org/>
- Two main components
 - YARN (for managing processing)
 - Distribute File System (HDFS) (for managing storage)

RDBMS Vs Hadoop

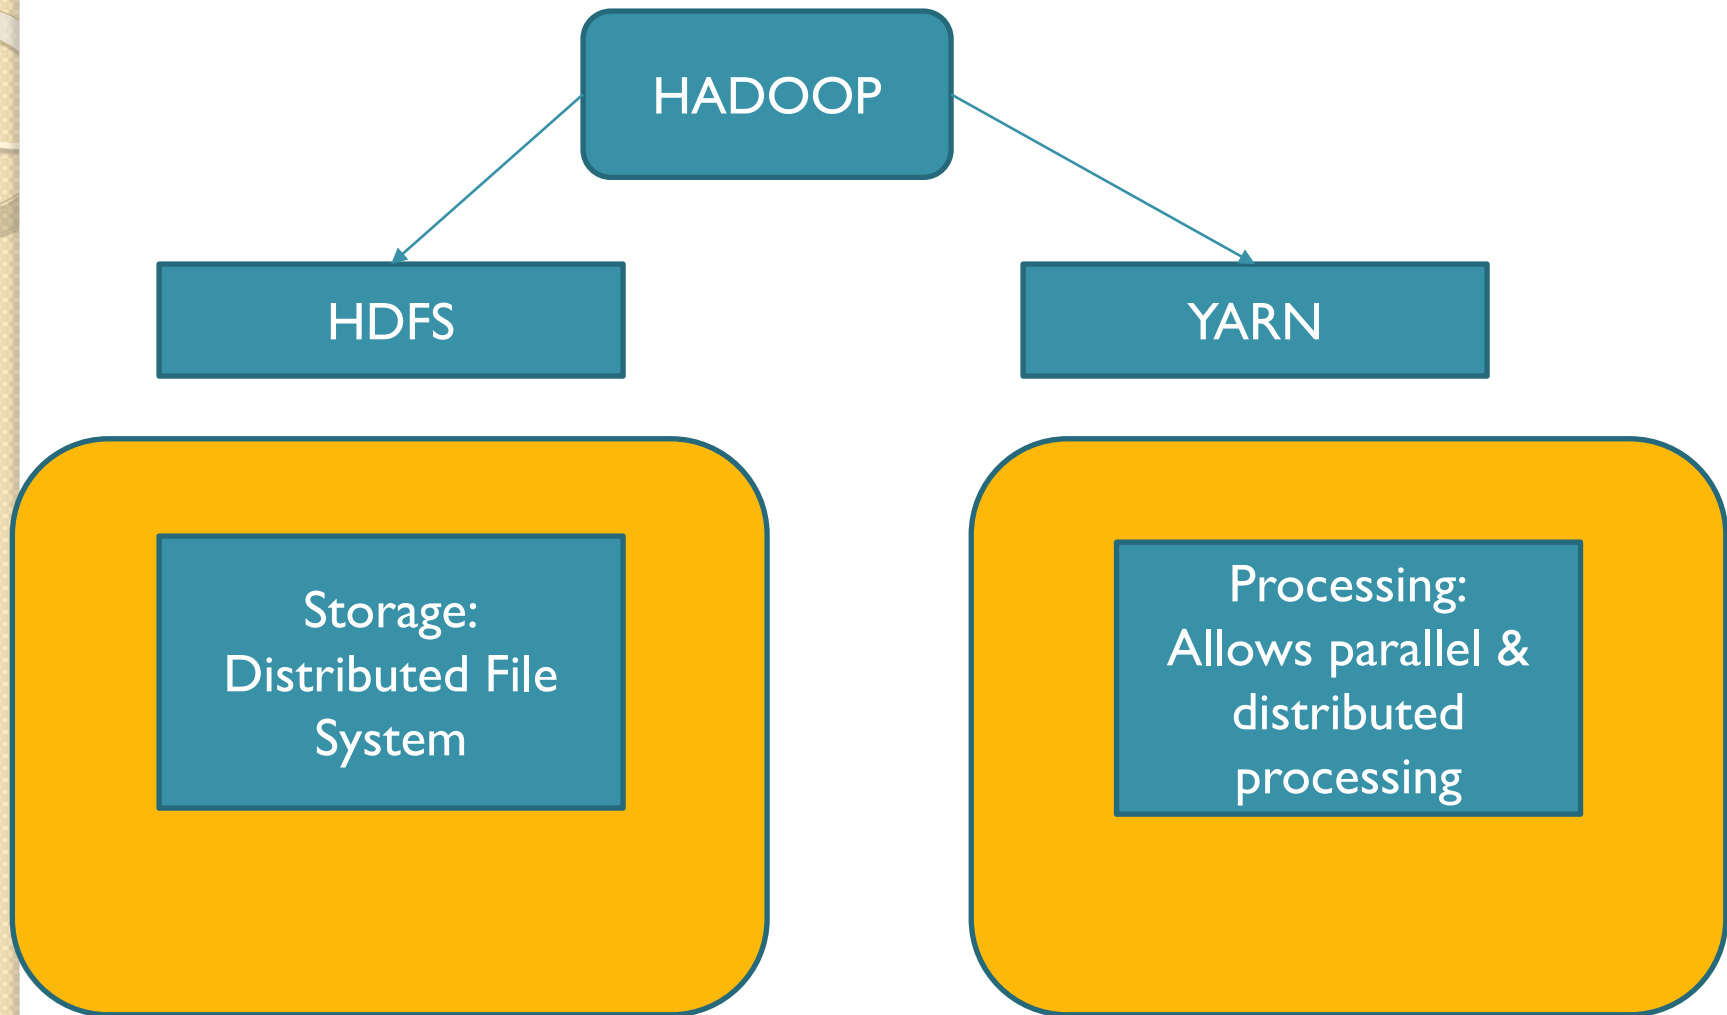
	RDBMS	Hadoop
Data Types	RDBMS relies on the structured data and the schema of the data is always known.	Any kind of data can be stored into Hadoop i.e. Be it structured, unstructured or semi-structured.
Processing	RDBMS provides limited processing capabilities.	Hadoop allows us to process the data which is distributed across the cluster in a parallel fashion.
Schema on Read Vs. Write	RDBMS is based on ' schema on write ' where schema validation is done before loading the data. CRUD supported	Hadoop follows the ' schema on read ' policy. Write once read many environment (WORM)
Read/Write Speed	In RDBMS, reads are fast because the schema of the data is already known.	The writes are fast in HDFS because no schema validation happens during HDFS write.
Cost	Typically licensed software, therefore, you have to pay for the software.	Hadoop is an open source framework.
Best Fit Use Case	RDBMS is used for OLTP (Online Transactional Processing) system. Is also used for OLAP\ data warehousing	Hadoop is used for Data discovery, ELT, data analytics or OLAP systems.



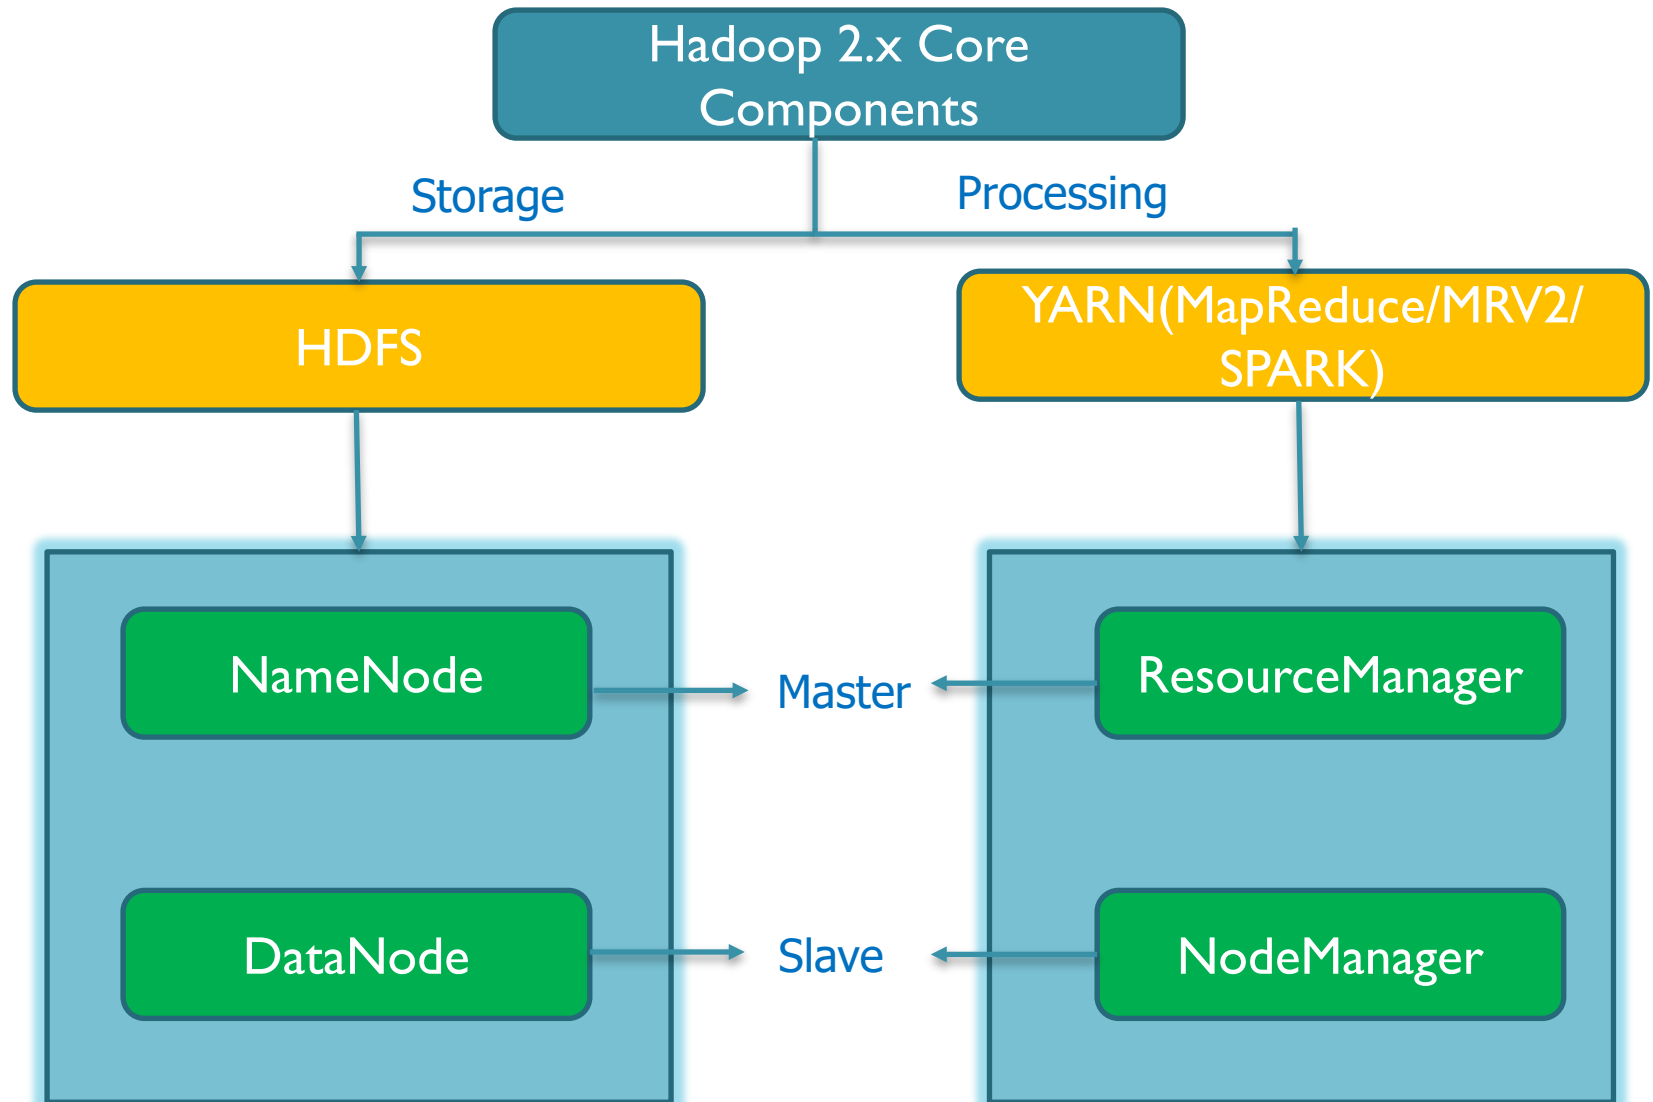
HADOOP Use Cases

- Data Warehouse Offload
 - Use Hadoop for long-term storage due to lower cost of storage and processing
 - Performing pre-integration ETL routines
 - Storing unprocessed data prior to being staged\integrated into a Data Warehouse
- Event Processing
 - Ingestion and processing of streaming data sources e.g. sensor data (e.g. RFID, CCTV cameras, temperature), log data , message data
 - Utilities include Storm, Flume, Spark Streaming, Kafka
- Advanced Analytics (Data Science)
 - Allow for machine learning\AI at scale

Hadoop-2.x




Hadoop2.x Main Daemons



NameNode and DataNode in HDFS

HDFS ARCHITECTURE



- NameNode

- Master Daemon
- Maintains and manages DataNodes
- Records Metadata e.g. location of blocks stored, the size of the files, permissions, hierarchy; for performance held in memory)
- Receives heartbeat and block report from all the DataNodes

- DataNodes

- Slave Daemons
- Stores actual data
- Serves read and write requests from the clients
- Sends a heartbeat to the NameNode

NameNode Metadata

Object	Block_id	Seq	Location	ACL	Checksum
/data/file.txt	Blk_00121	1	[DN1,DN2,DN3]	-rwxrwxrwx	8708b09....
/data/file.txt	Blk_00122	2	[DN2,DN3,DN4]	-rwxrwxrwx	cd786a87..
/data/file.txt	Blk_00123	3	[DN2,DN4,DN5]	-rwxrwxrwx	cd786a87..

Conceptual Representation of In_Memory Metadata


HDFS Blocks

- Each File is stored on HDFS Blocks
- The default size of each block is 128mb in Apache Hadoop 2.x (64mb in Apache Hadoop 1.x)
- Let us say we have a file 376mb it will be split into chunks as follows:



Blocks are Distributed


- If cluster contains more than one node, blocks are distributed



Note: Each Block is also replicated in the DataNodes (Not shown in the diagram)

HADOOP Block Replication

Data File (500mb)



Files broken up
into 4 blocks (chunks)
With a Replication Factor(3)


Data Node 1



Data Node 2



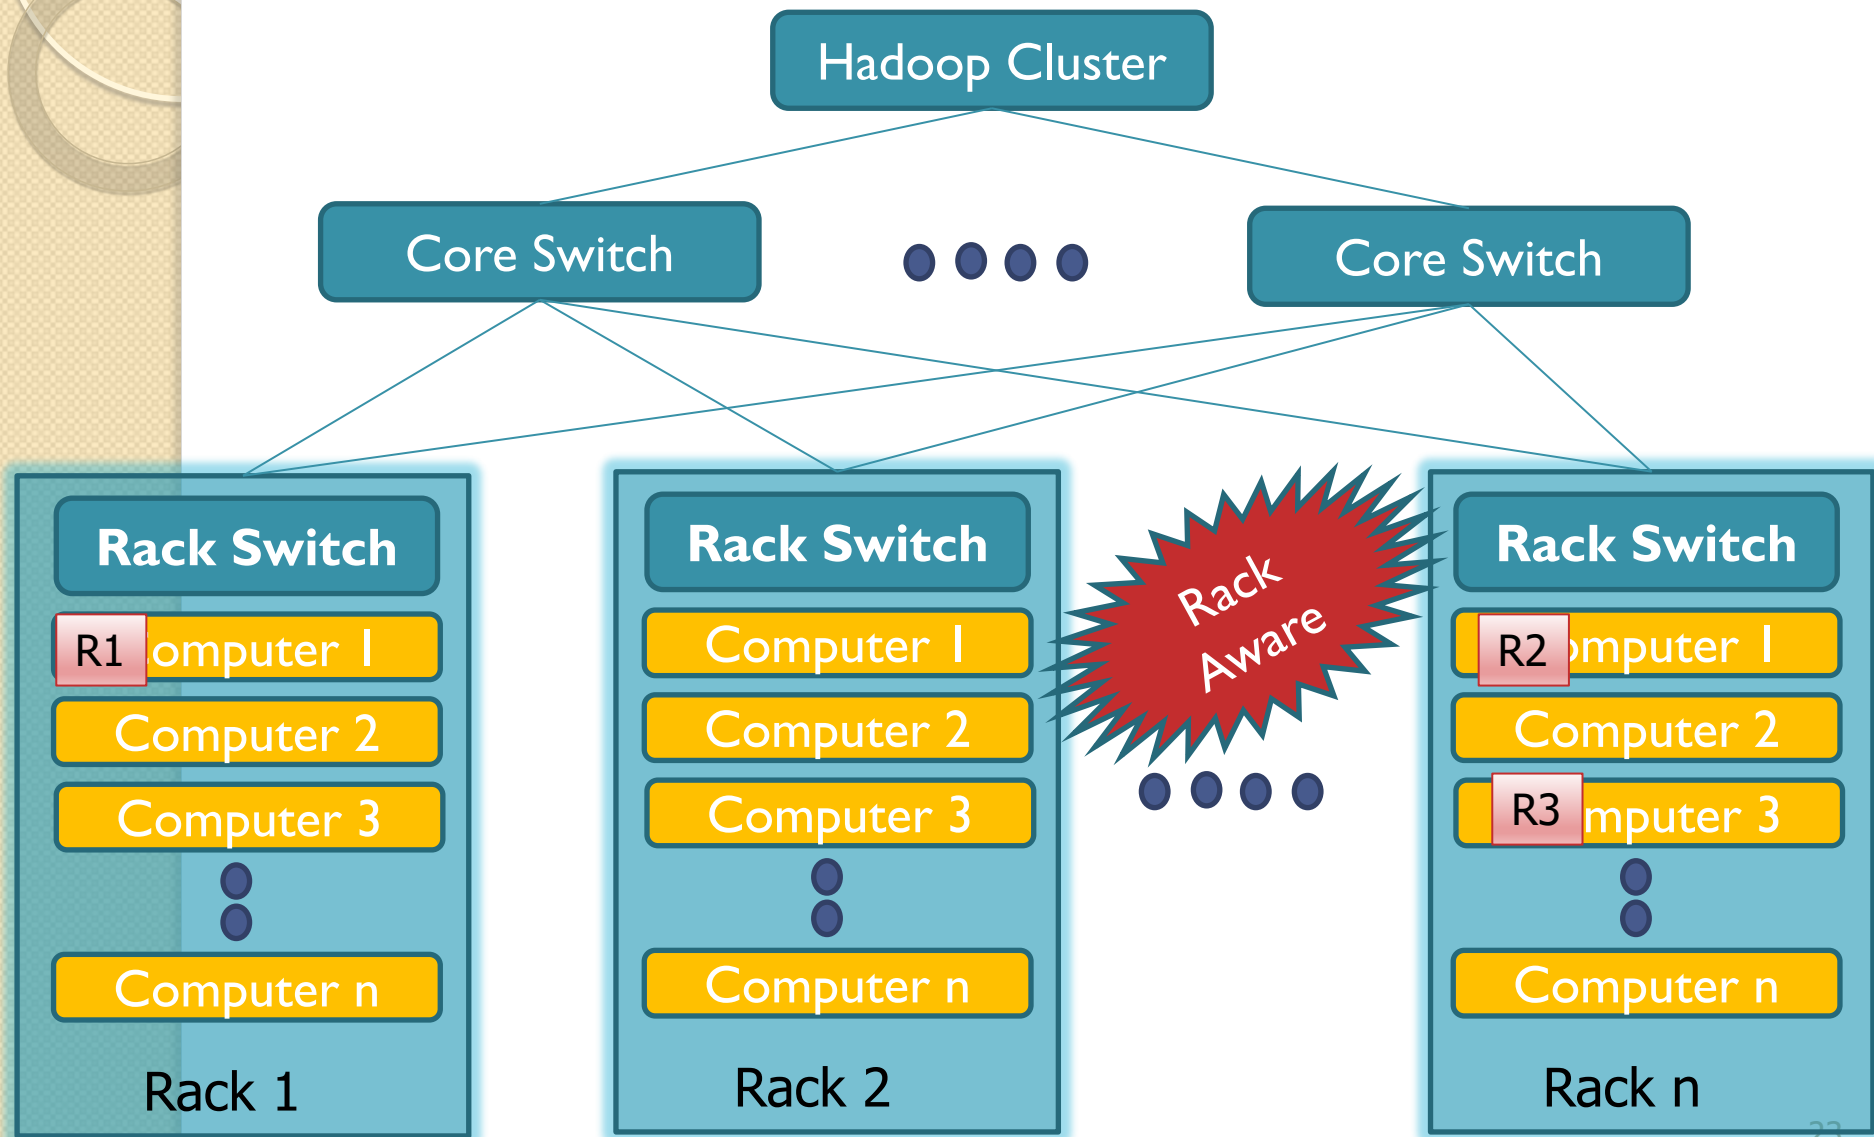
Data Node 3



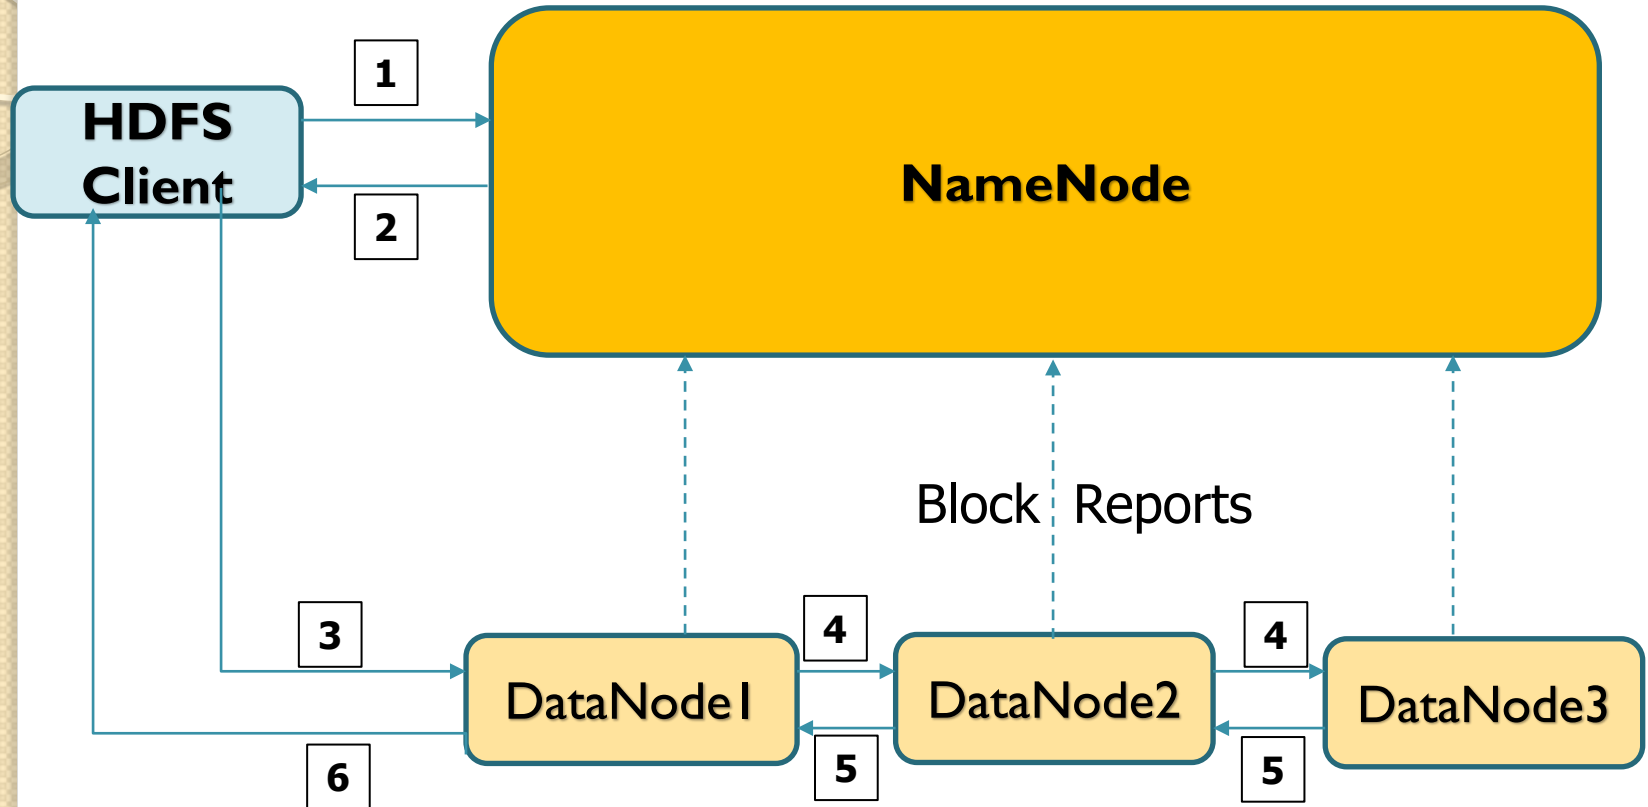
Data Node 4



Hadoop Cluster Architecture

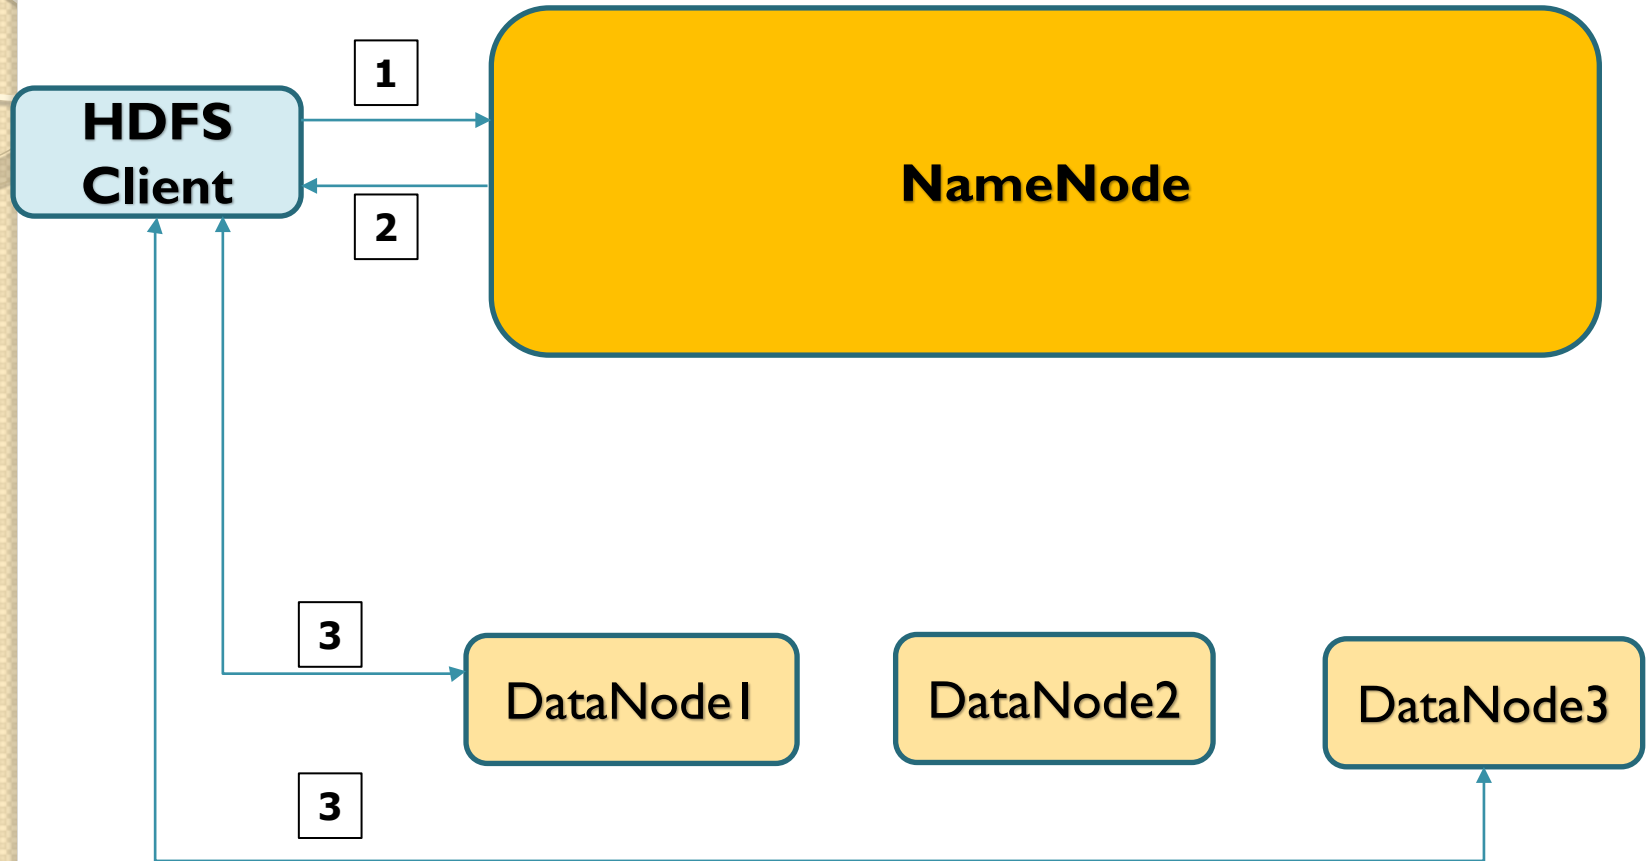


A Write Operation




“Block Replication Pipeline”

A Read Operation



Data Node Failure



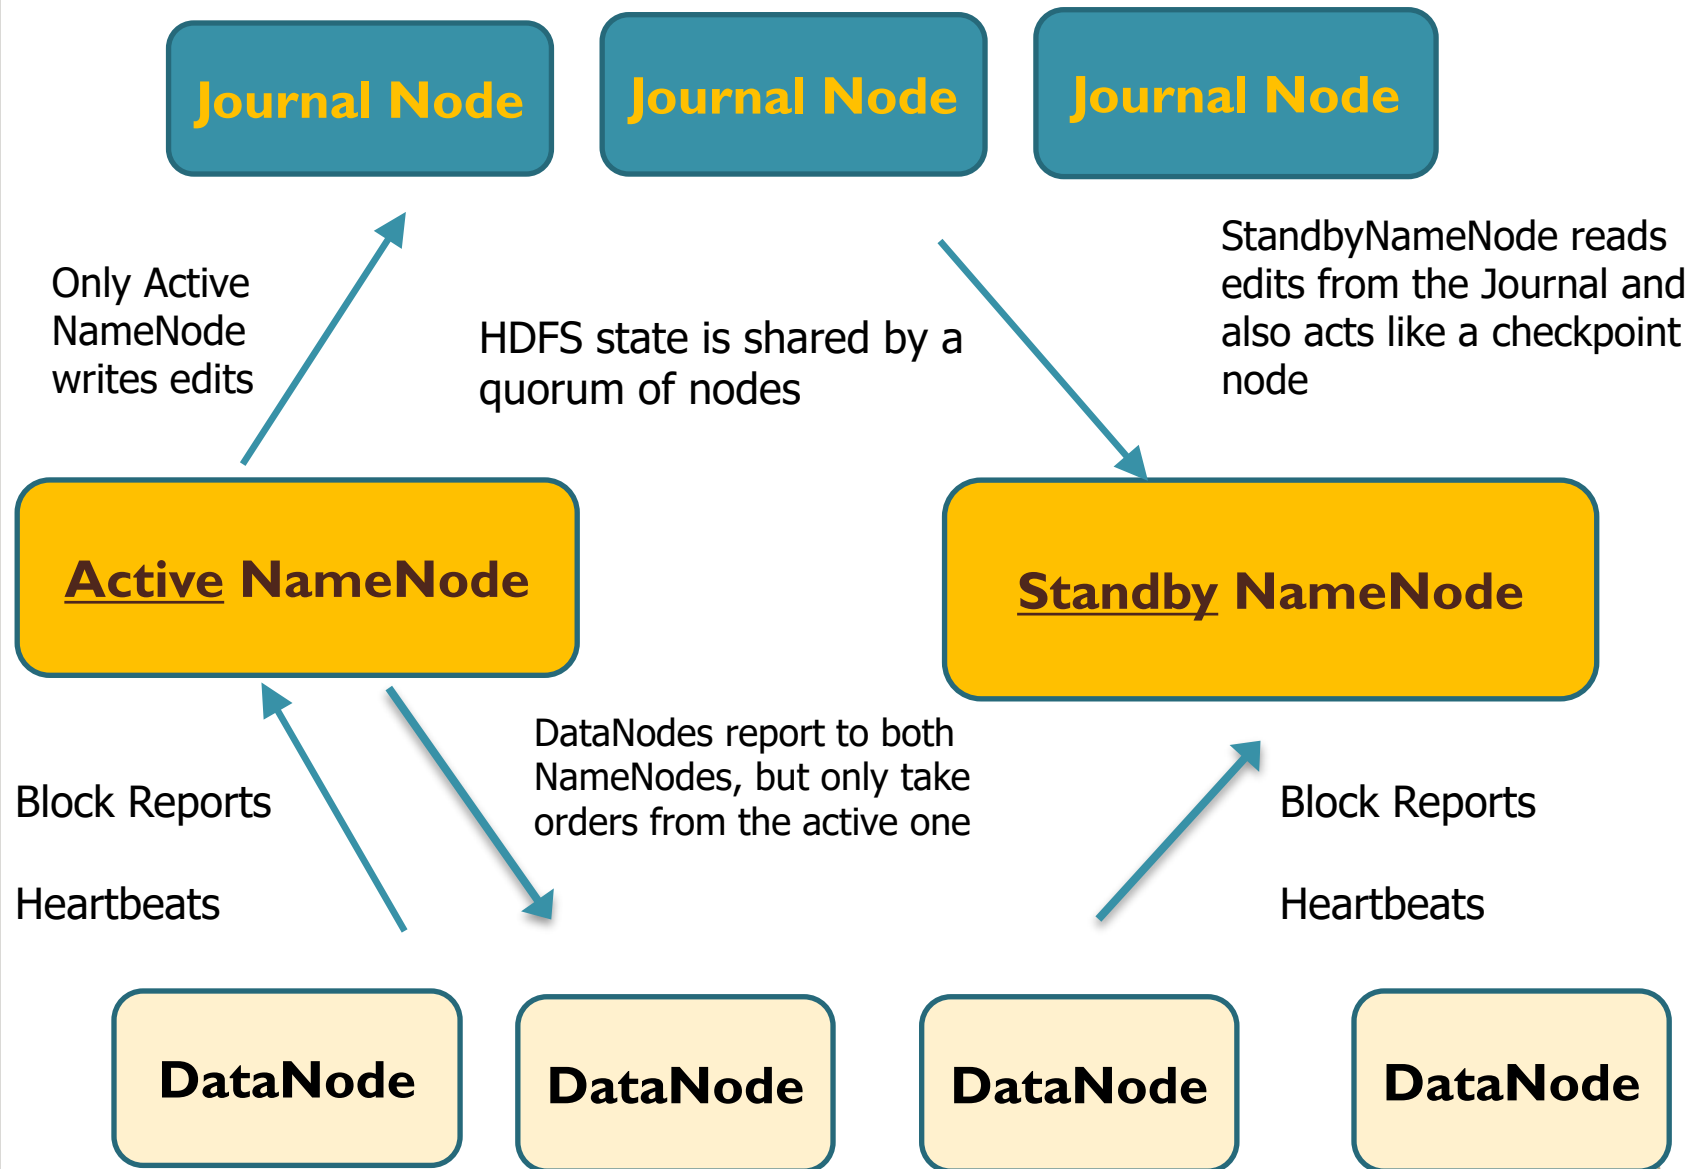
- NameNode detects no heartbeat from DataNode
- Marked as **dead** after specific period of time
- NameNodes replicates blocks to another Datanode (using the replicas created earlier)
- **Note** NameNode has a list of all the blocks on the Datanode(known as a block report)

Secondary NameNode(2NN) - Optional

- Performs periodic **checkpoints** that evaluate the status of the NameNode
- Main aim is to **improve NameNode restarts**
- NameNode keeps Metadata in memory. MetaData contained on disk in
 - **fsimage_**
 - An image snapshot of the HDFS file state when the Name-node was started
 - **edit_***
 - A series of modifications made to HDFS during the running of the NameNode (like a transaction log in RDB) or REDO log in Oracle
- Secondary NameNode
 - Periodically downloads these files from the Namenode
 - Merges them and loads an image back to the NameNode

***Secondary NameNode not an active failover node**

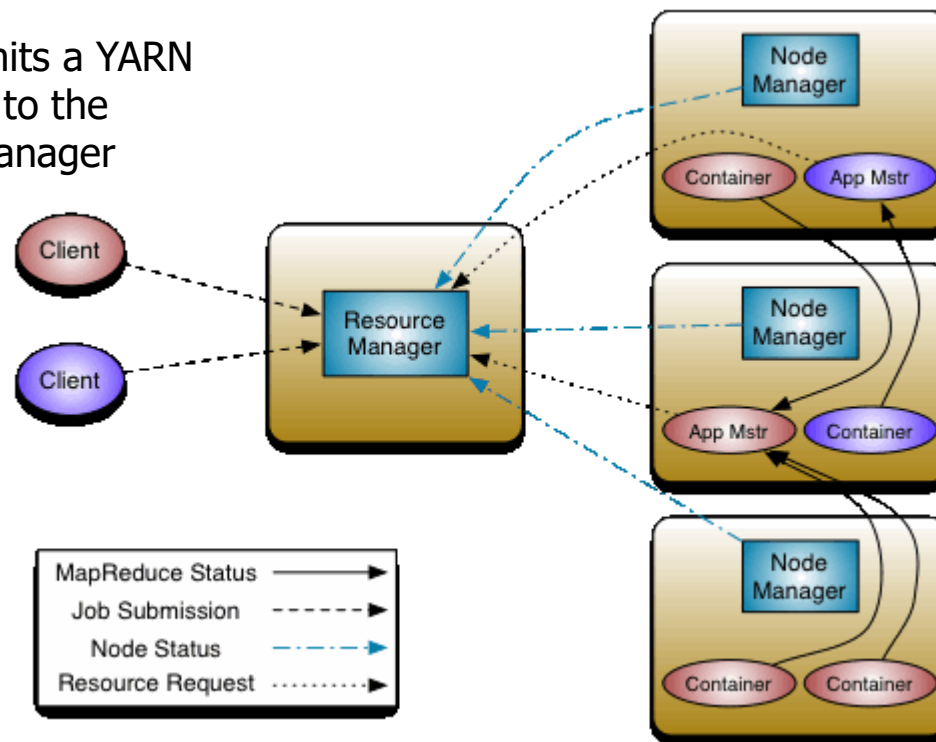
Standby NameNode for HA -Optional



YARN (Yet Another Resource Negotiator)

The ResourceManager assigns an Application Master (a container) for the Application

Client submits a YARN application to the ResourceManager



Tasks for an application report their progress to the ApplicationMaster for the application

Map/Reduce

❑ MapReduce

- ❑ Distributed computation framework
- ❑ Optimized for batch processing
- ❑ Typical I/O profile:



Count words in docs – The “Hello World” Of Map Reduce

- Input consists of (filename, file-contents) pairs

- `map(key=filename val=file-contents):`
 - For each word `w` in file-contents, emit (`w`, “1”)

- `reduce(key=word, values=uniq_counts):`

Sum all “1”s in values list

`sum = 0`

For each value in values:

`Sum = sum + value`


Emit result “(word, sum)”

Word Count, Illustrated

- `map(key=filename val=file-contents):`
 - For each word `w` in `file-contents`, emit `(w, "1")`
- `reduce(key=word, values=uniq_counts):`
 - Sum all "1"s in values list
 - `sum = 0`
 - For each value in values:
 - `Sum = sum + value`
 - Emit result `"(word, sum)"`


see bob throw
see spot run

map



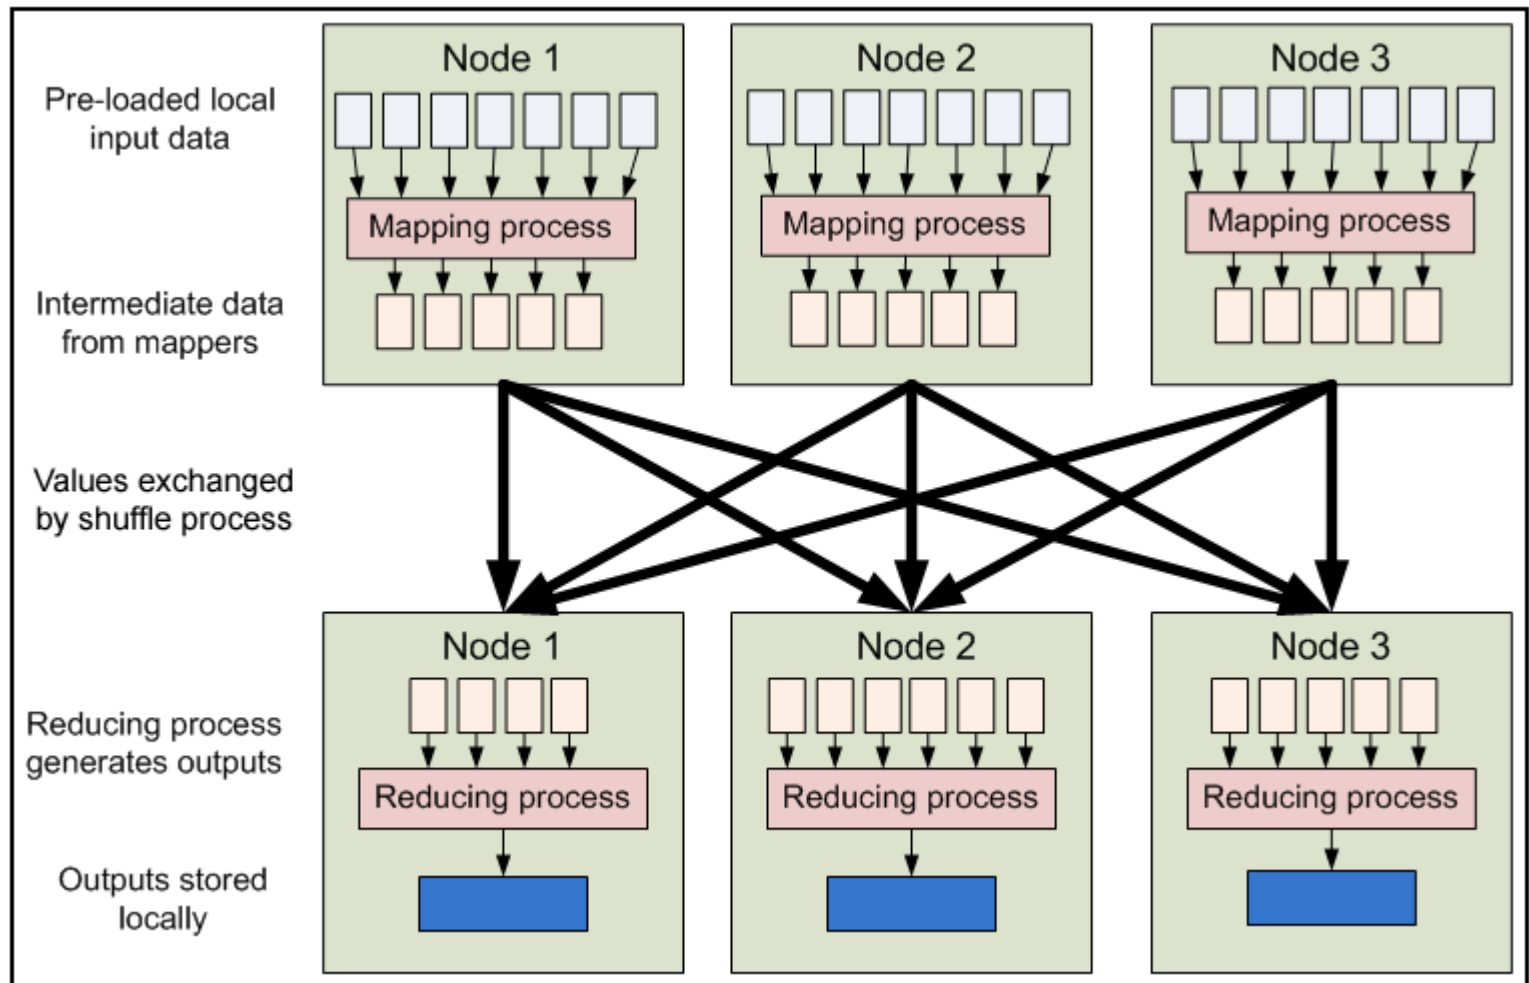
see	1
bob	1
run	1
see	1
spot	1
throw	1

reduce

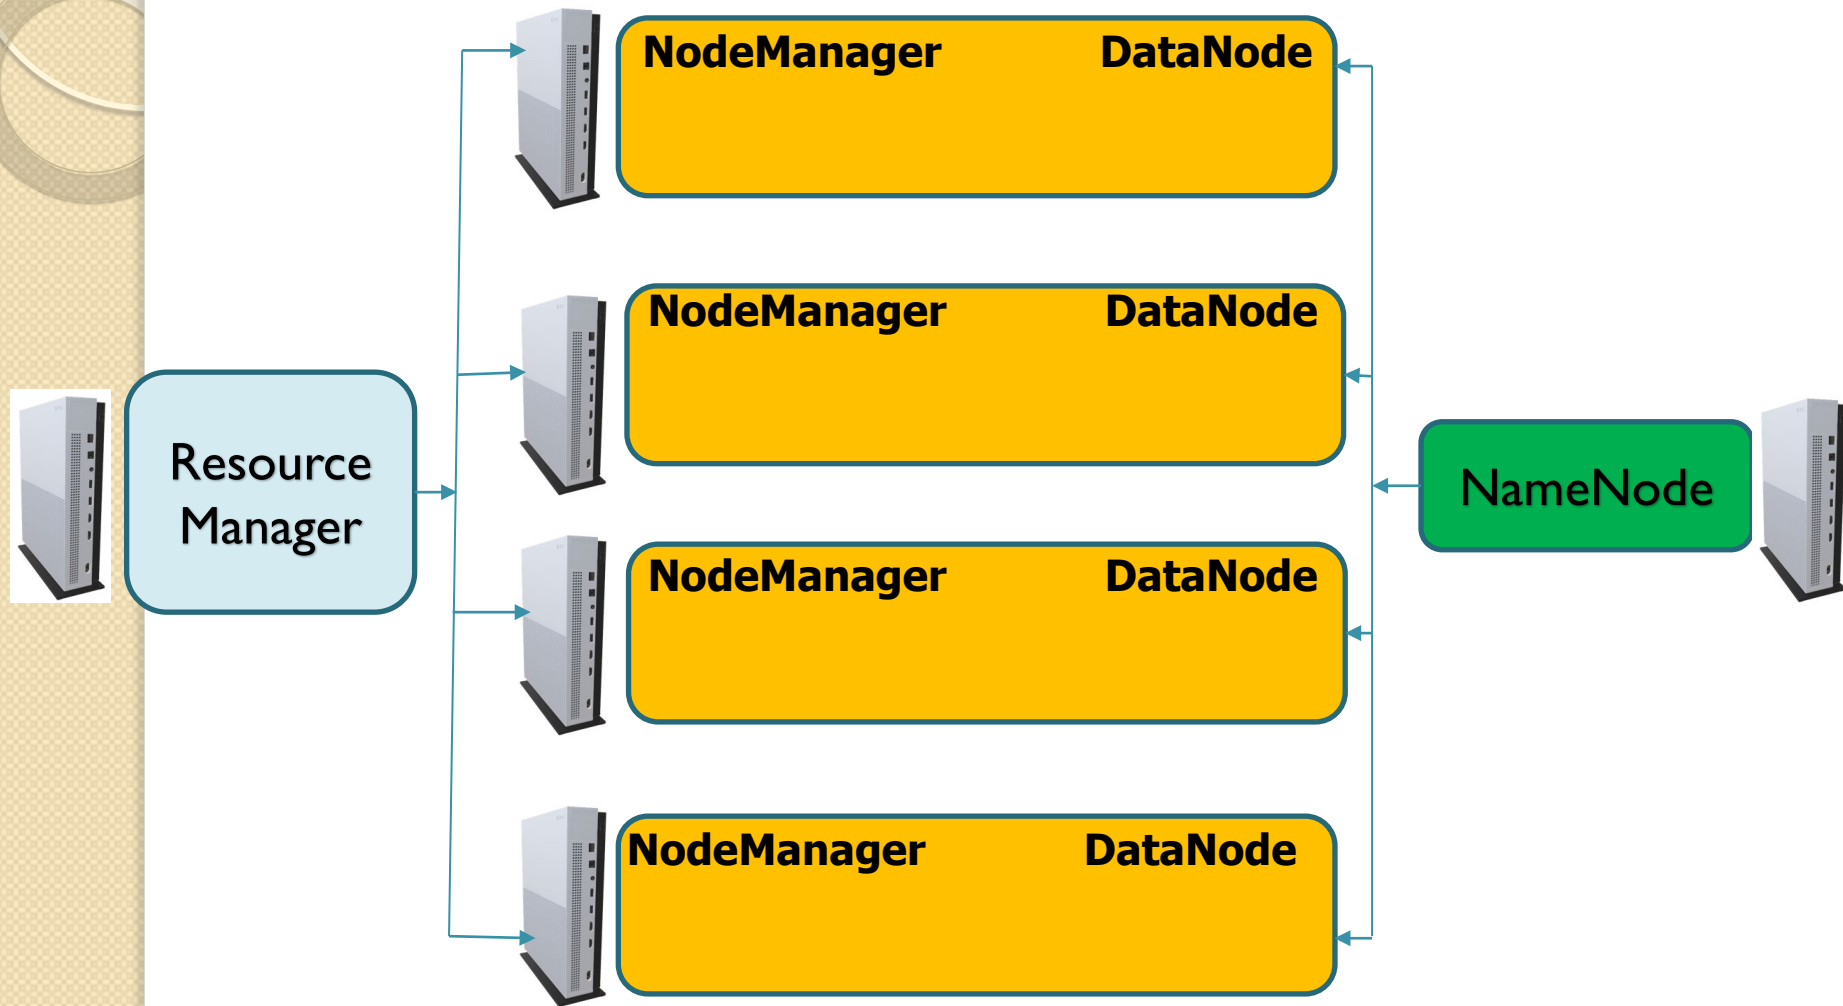


bob	1
run	1
see	2
spot	1
throw	1

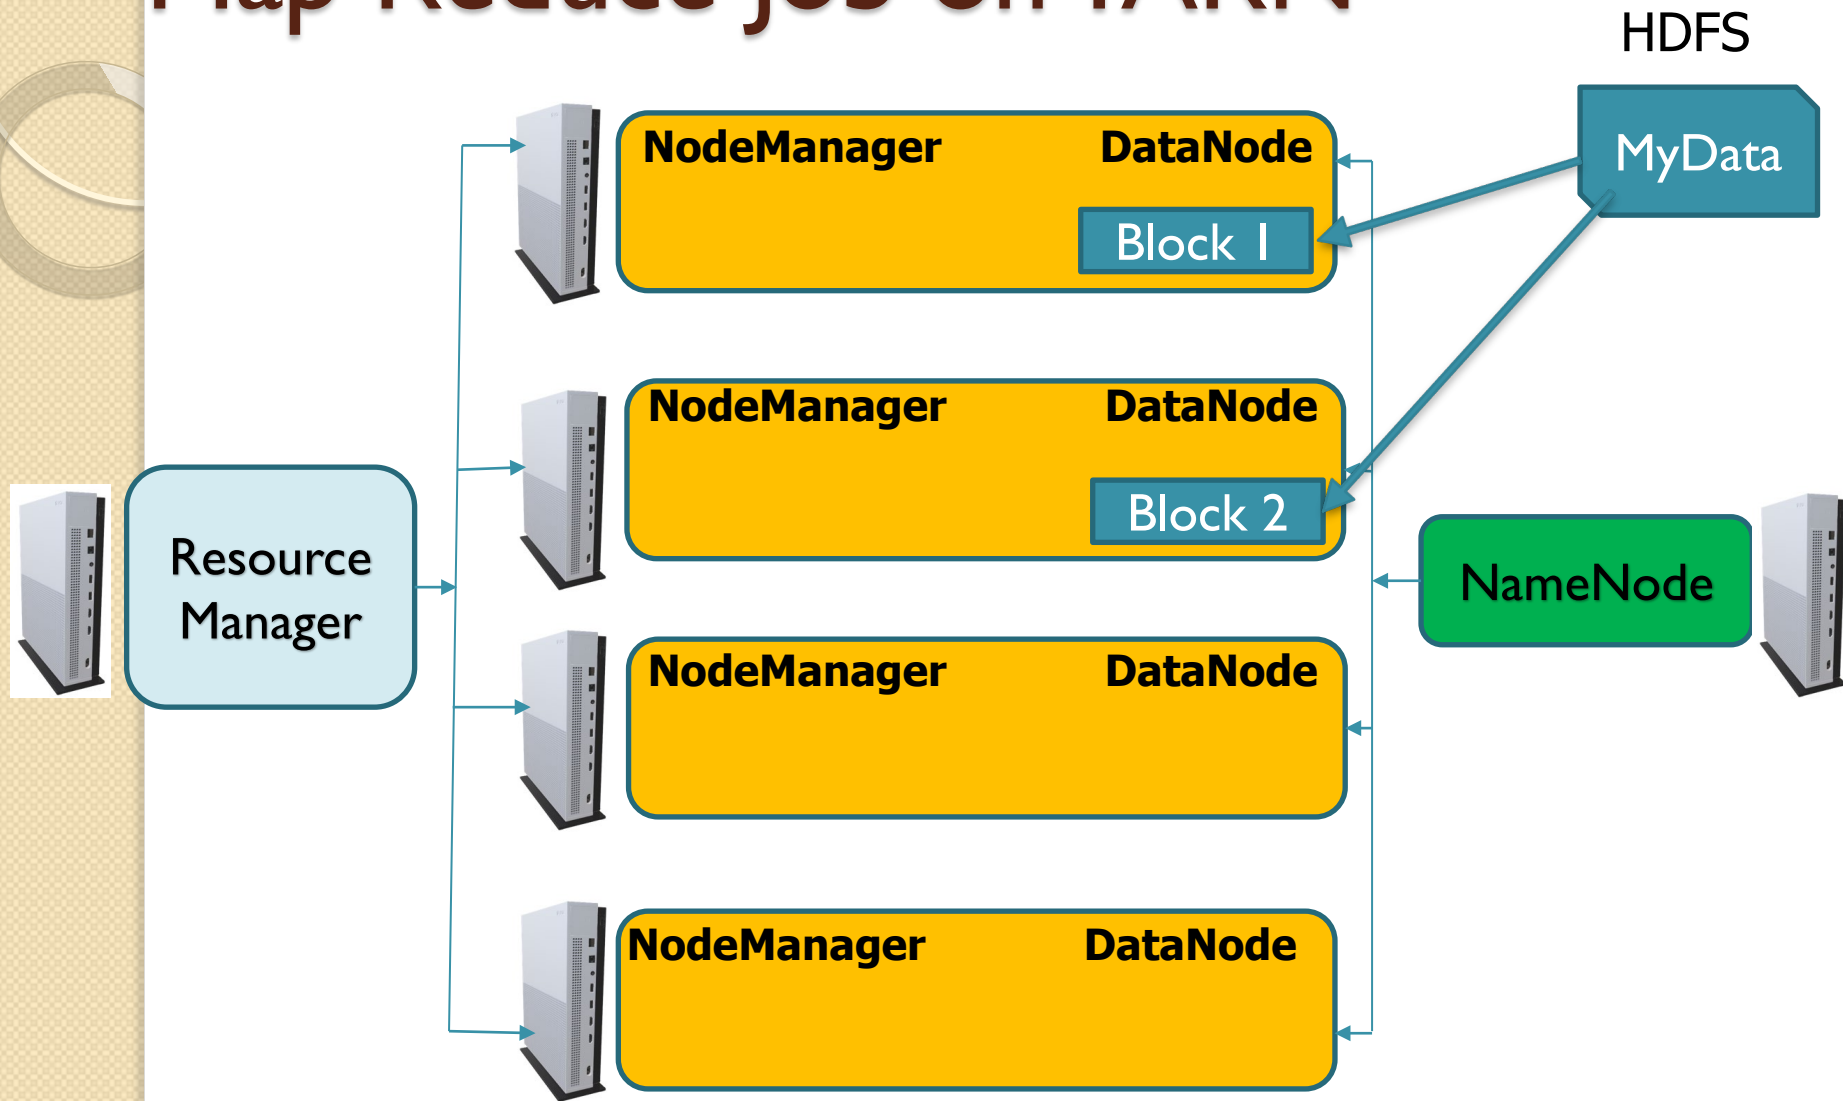
Mapping and Reducing tasks run on data nodes containing individual records of data



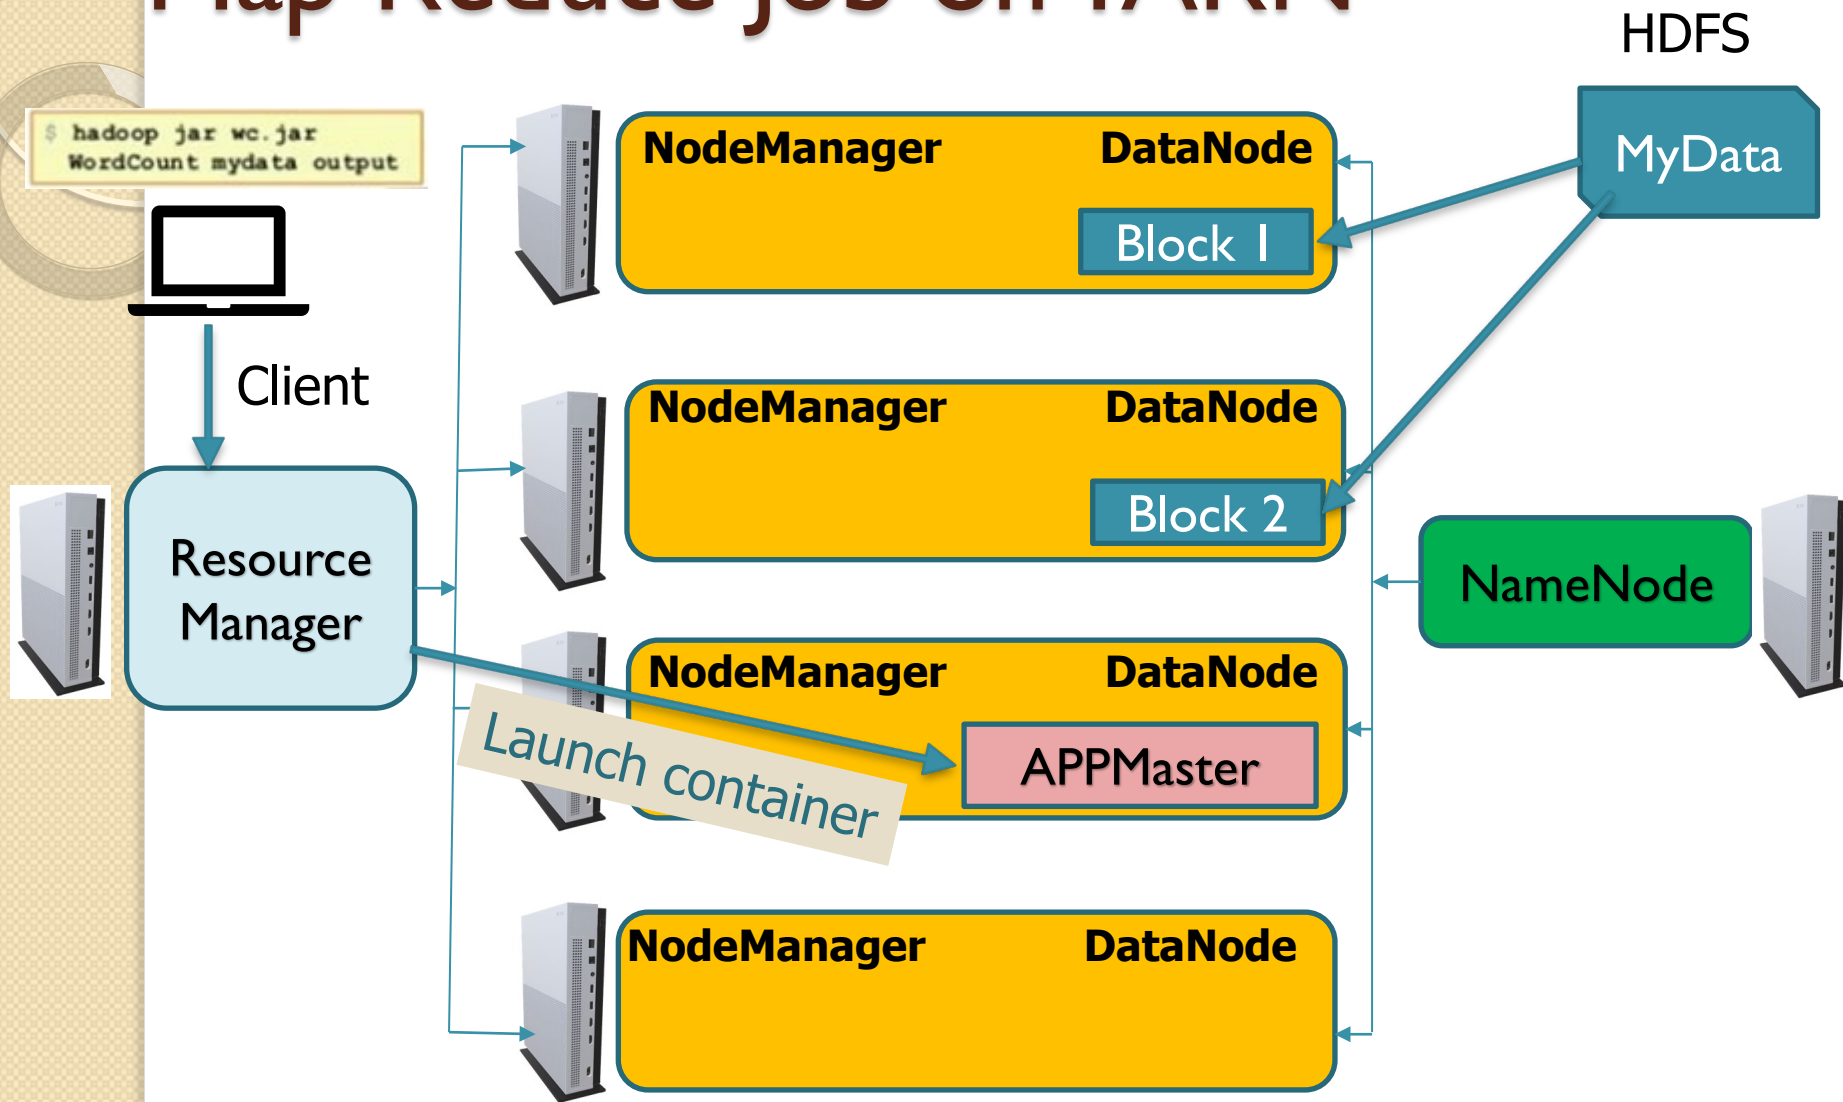
Map Reduce Job on YARN



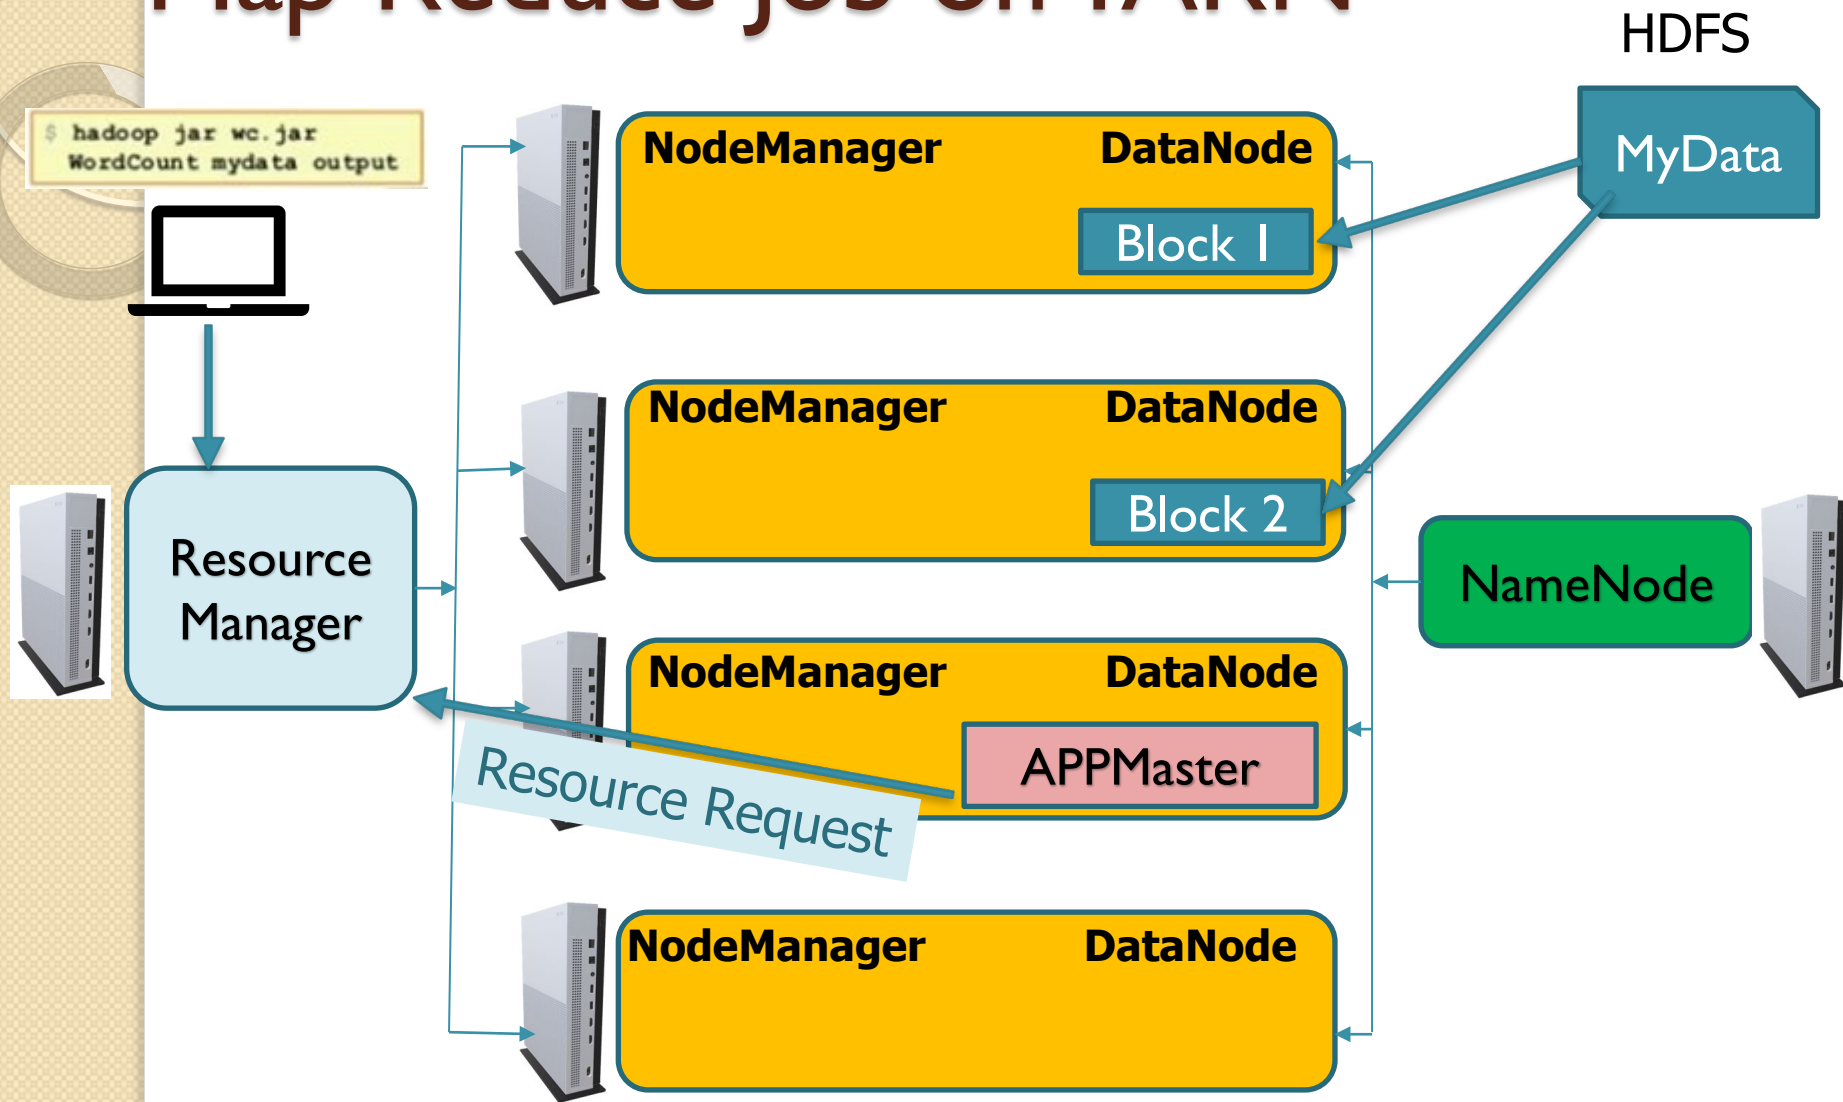
Map Reduce Job on YARN



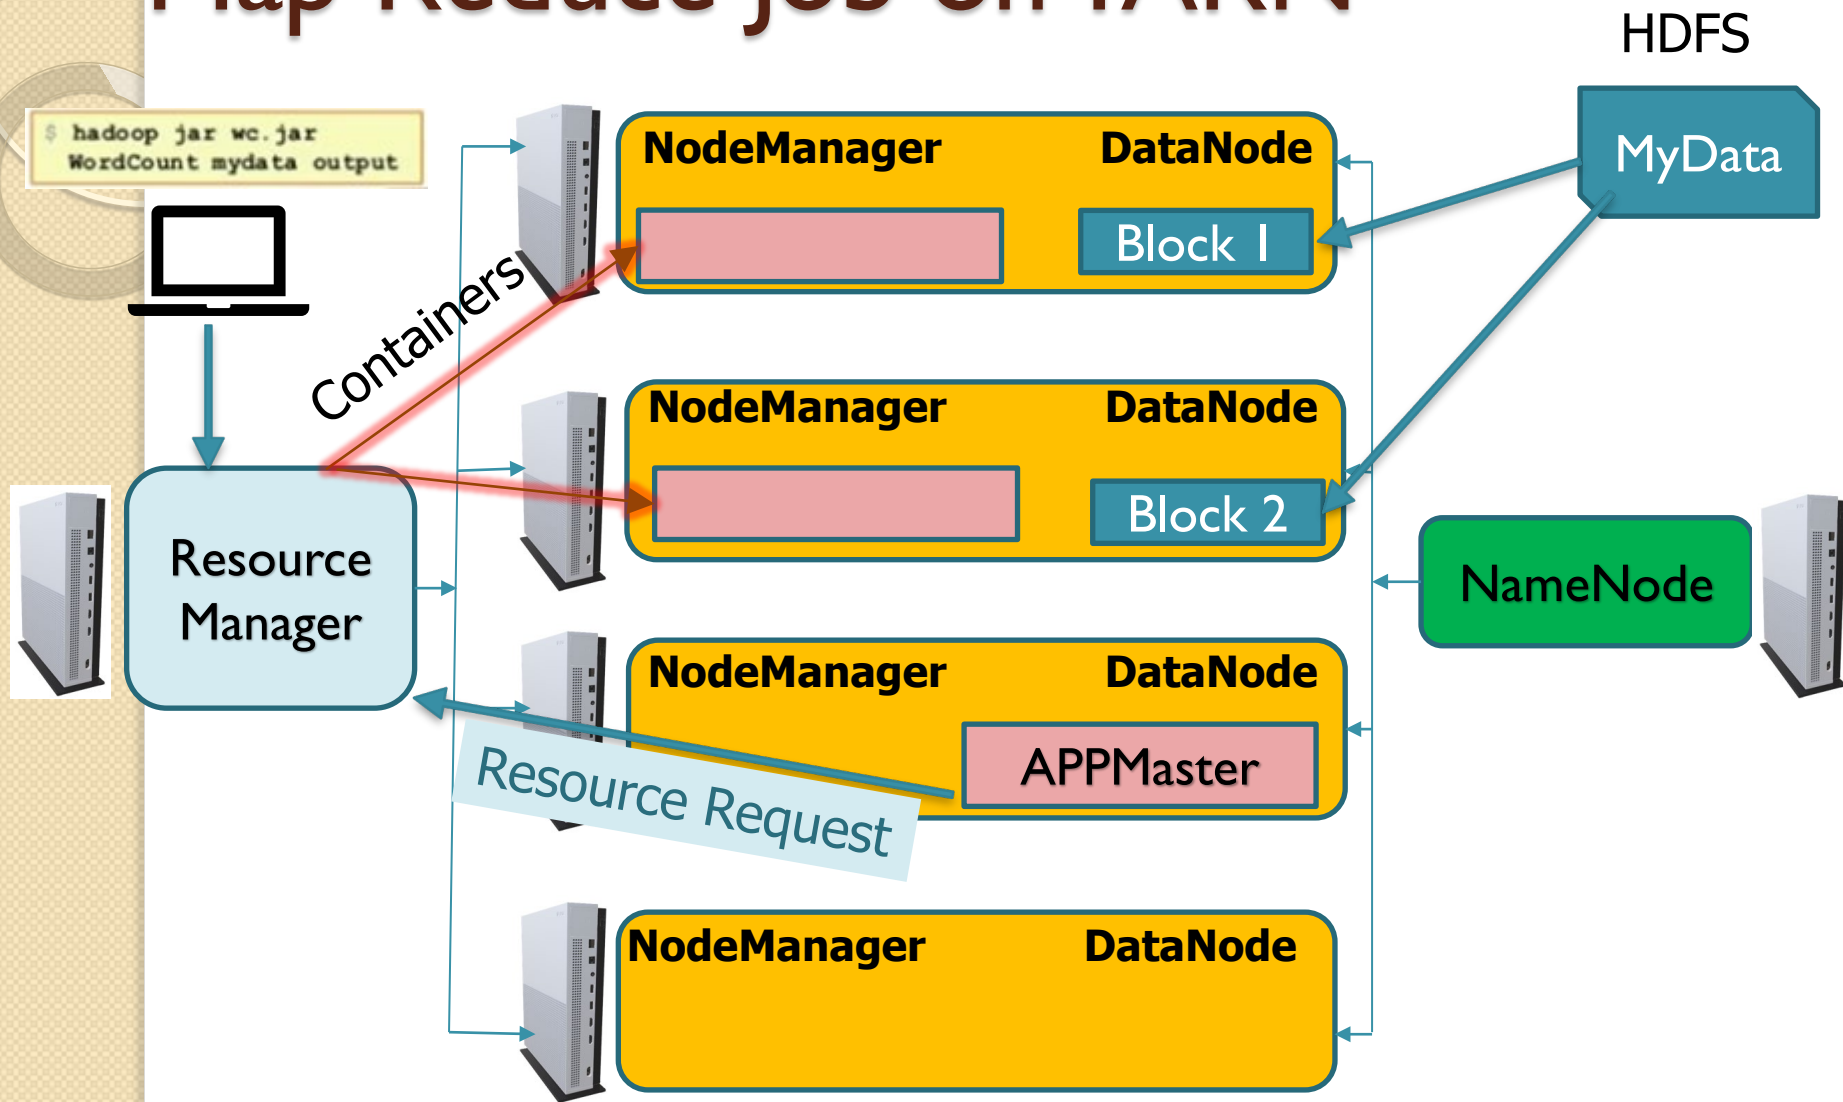
Map Reduce Job on YARN



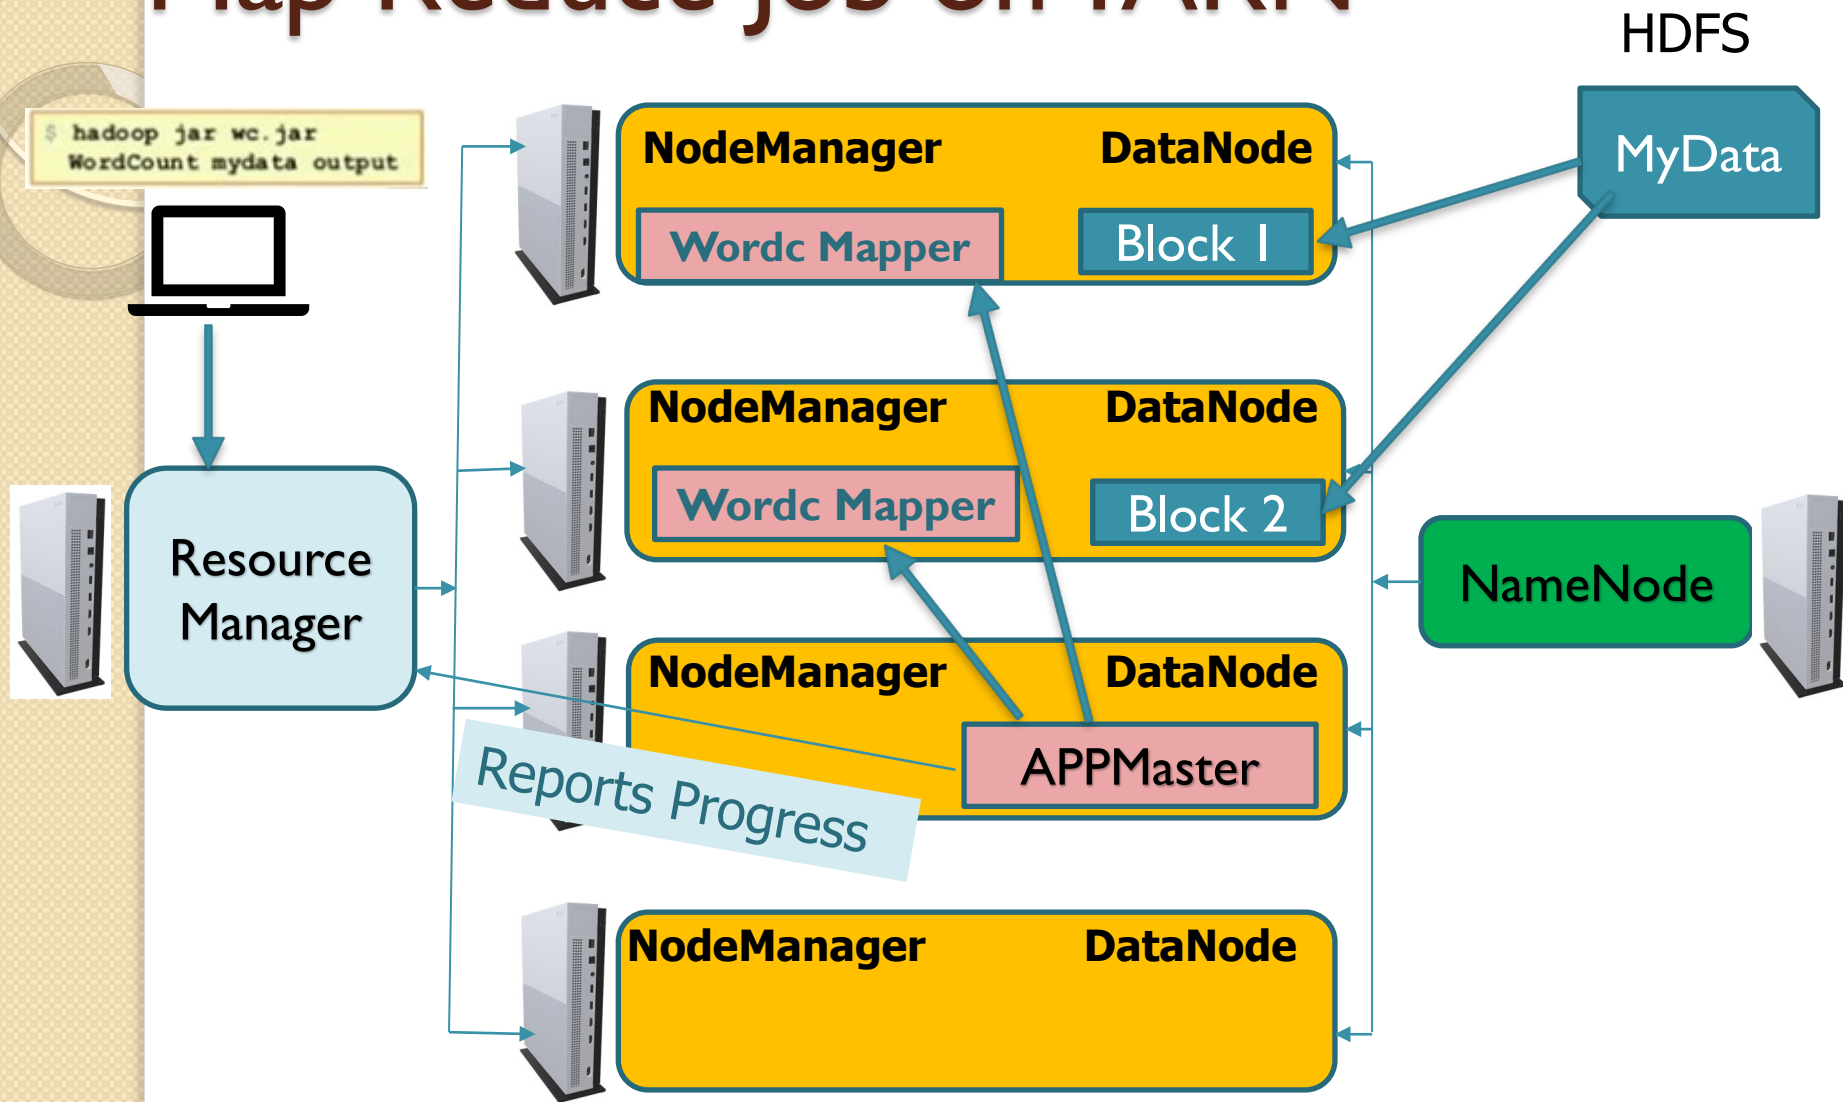
Map Reduce Job on YARN



Map Reduce Job on YARN

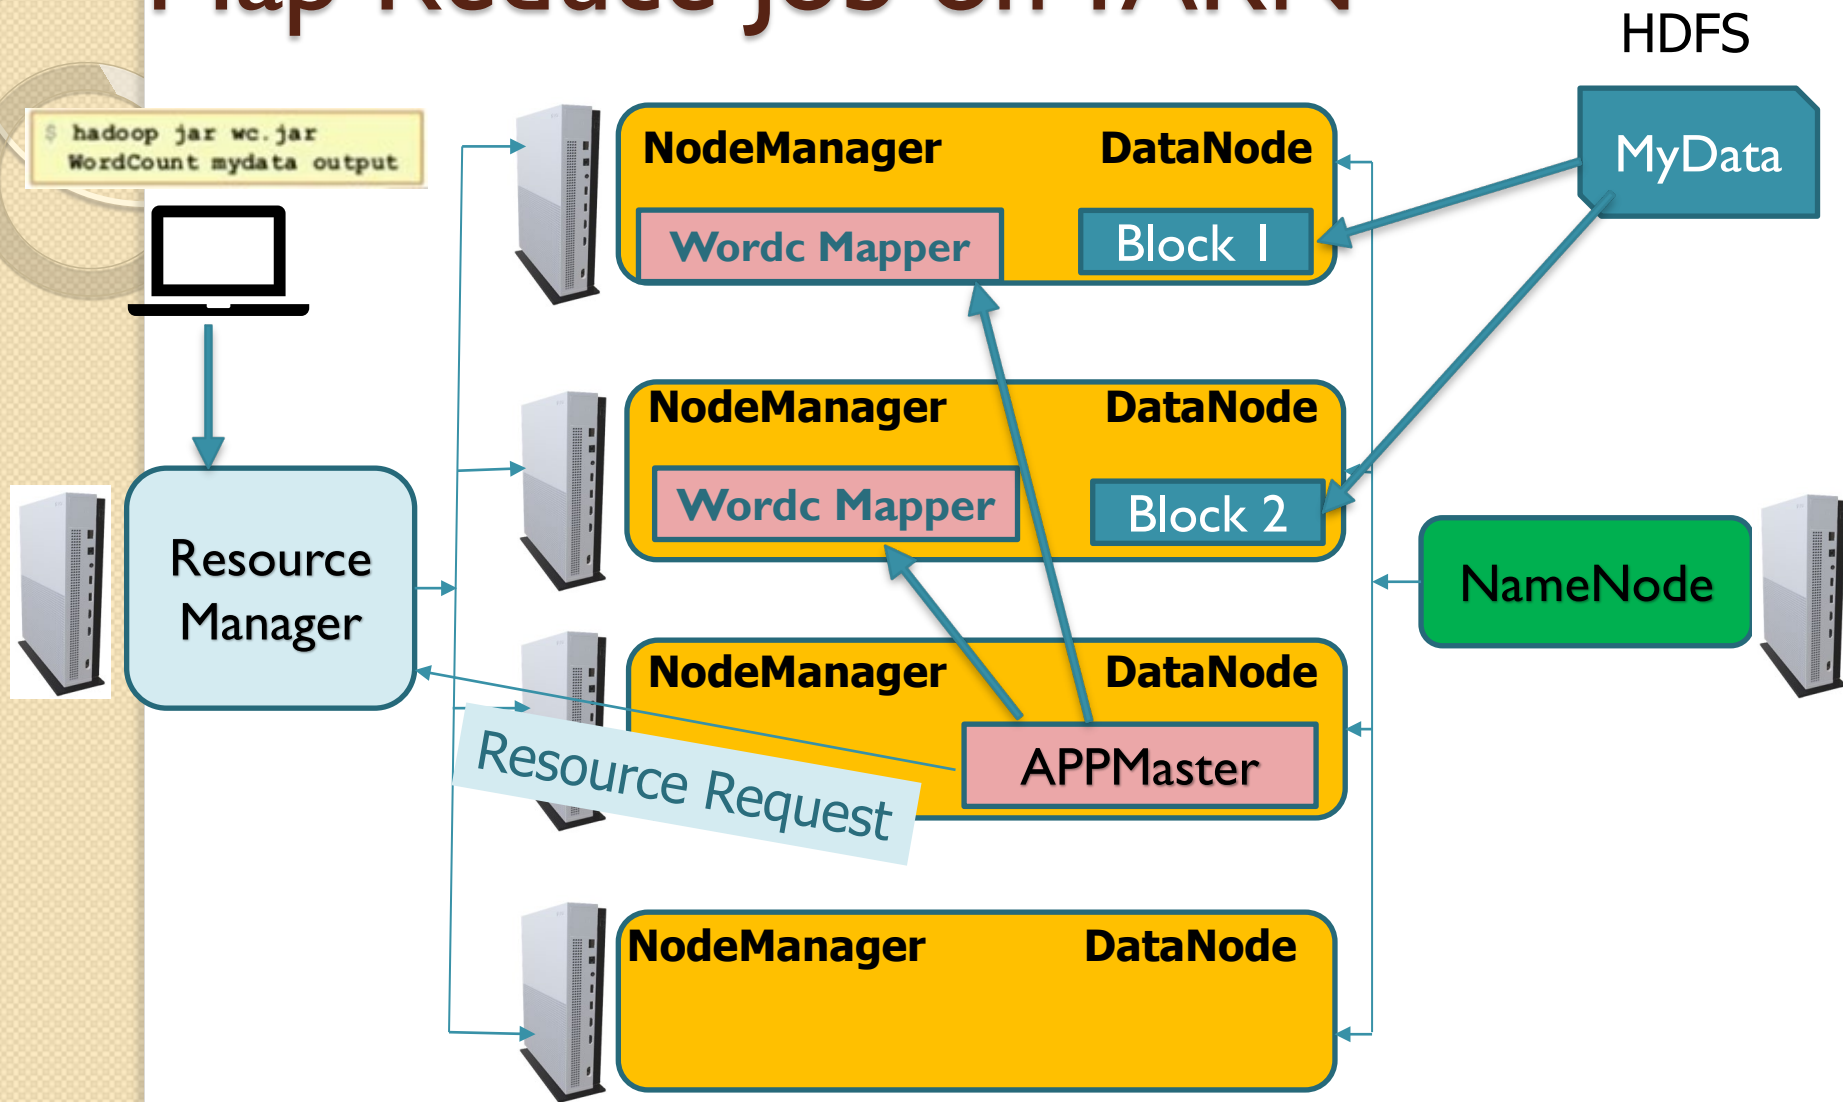


Map Reduce Job on YARN

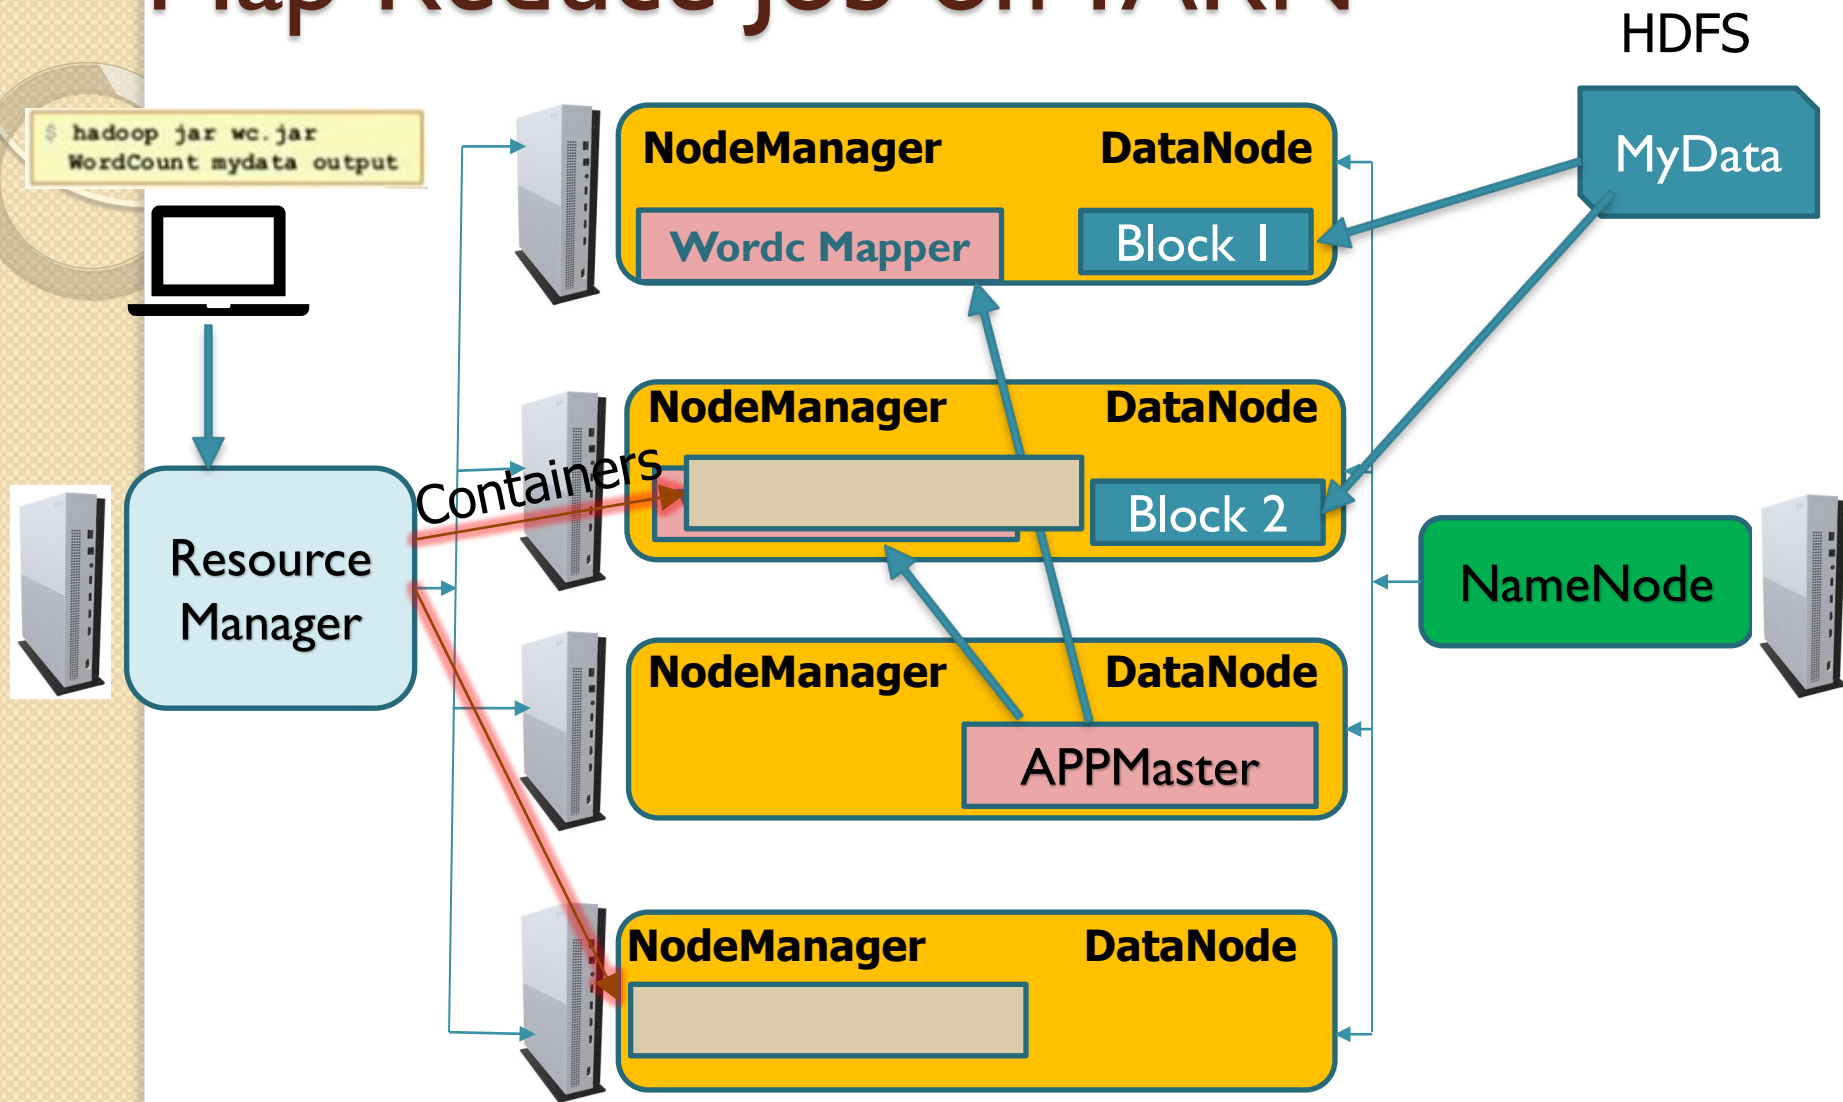


Note: Reducer Containers required

Map Reduce Job on YARN



Map Reduce Job on YARN



Map Reduce Job on YARN

