



# NOSQL Databases



Part 3 – NoSQL Databases


# NoSQL Implementation Categories

---

- ▶ Key-Value Databases
- ▶ Document Databases
- ▶ Column-Family stores
- ▶ Graph Databases

# Key Value Stores

- ▶ Simplest of the NoSQL storage mechanisms
  - ▶ Focus on ability to store and retrieve data rather than data structure
  - ▶ Get a value for the key; put a value for a key; delete a key from the data store
- ▶ Essentially a large hash table
- ▶ Key/Value Pair
  - ▶ **Key** is unique to identify the data and value;
    - ▶ like an Oracle ROWID
    - ▶ All data access via the Key (mostly!)
  - ▶ **Value** is often opaque; just stores data without caring what was stored. The value can be a blob, text, JSON, XML etc.
    - ▶ To store the value does not have structure
    - ▶ Responsibility of the application to understand what was stored; DBMS does not care!
- ▶ Can work in a distributed partitioned replicated environment



ORACLE®

BERKELEY DB



# Key Values and Associative Array

---

## ► An Associative Array

- Like an Ordered array but fewer constraints on keys and values

1	True
2	True
3	False
4	True
5	False

'Pi'	3.14
'CapitalFrance'	'Paris'
17234	34468
'Foo'	'Bar'
'Start_Value'	1

## ► Key-Values build on this idea

- Many key-value data stores keep persistent copies of data on long-term storage, such as hard drives or flash devices (e.g. Riak)
- Some key-value data stores only keep data in memory. (e.g,MemcacheD)



# Key And Nothing But The Key!

---

- ▶ Key Design critical to ensure they are easily managed
- ▶ RDBMS Keys a sequence of values e.g. I2234, I2235, I2236
  - ▶ Keys are not necessarily related to the purpose of information stored
- ▶ A prefix could be added: custI2234, custI2235, custI2236
  - ▶ Indicates that the keys refer to values about customers
- ▶ A more complex naming pattern e.g.
  - ▶ cust: I2234 :firstName
  - ▶ cust: I2234 :lastName
  - ▶ cust: I2234 :shippingAddress
- ▶ Developers must create a key naming convention
  - ▶ Keys too long – RAM considerations- kv databases are memory intensive
  - ▶ Keys too short – could lead to conflicts in key names



# The Value

---

- ▶ **A Value is an object associated with a key**
  - ▶ Can be simple integers, floating-point numbers, strings of characters, BLOBs, semi-structured constructs such as JSON objects, images etc.
  - ▶ Most key-value databases will have a limit on the size of a value ( multiple megabytes to a couple of kilobytes of data)
- ▶ **Key-value databases typically treat values as atomic opaque units**
  - ▶ Exceptions to this are key-value databases that provide text-search capabilities; this feature is not a standard part of key-value databases
    - ▶ E.g. Riak allows searching using Apache Solr
    - ▶ DynamoDB recommends queries using the key. However, allow search scans of the items (their value!)



# Key Values and Namespaces

## ► Namespace:

- A collection of key-value pairs.
- A Set; duplicate keys are not allowed.
  - Note It is possible to have the same value assigned to multiple keys, but keys can be used only once in a namespace.
- A namespace could be an ENTIRE key-value database OR
- A namespace can be at Bucket Level

Database

Bucket 1	
'Foo1'	'Bar'
'Foo2'	'Bar2'
'Foo3'	'Bar7'

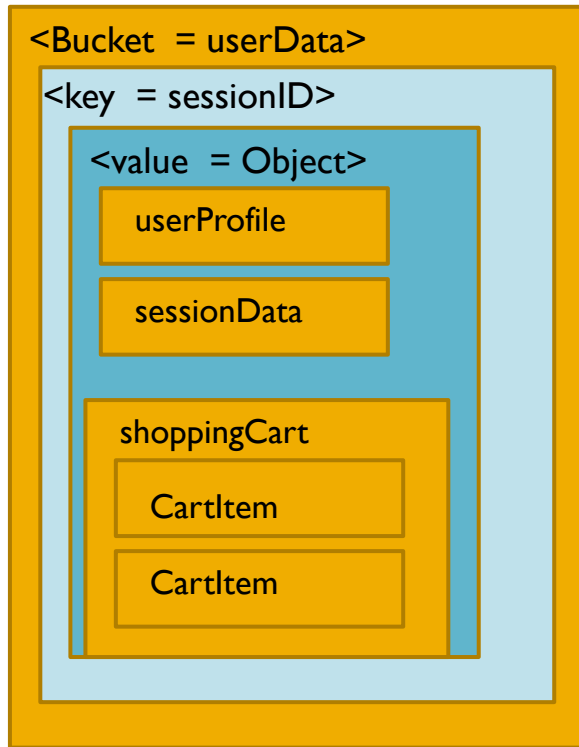
Bucket 2	
'Foo1'	'Baz'
'Foo4'	'Baz3'
'Foo6'	'Baz2'

Bucket 3	
'Foo1'	'Bar7'
'Foo4'	'Baz3'
'Foo7'	'Baz9'

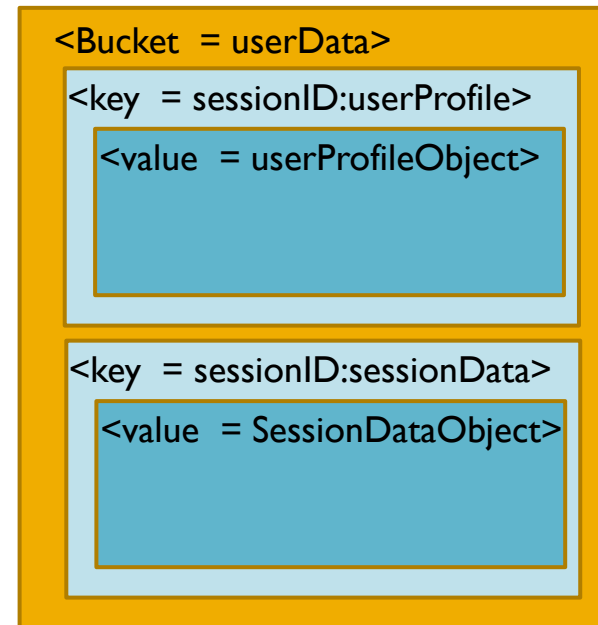
- The essential characteristic of a **namespace** is that it is a collection of key-value pairs that **has no duplicate keys**.



# Riak Bucket



Storing all data in a single bucket



Change the key design to segment the data in a single bucket

## Key Design is important in KV stores



# DynamoDB - Tables, Items, and Attributes

## A Table

Primary Key  
- Partition Key  
- Sort Key

Primary Key  
(Partition Key)

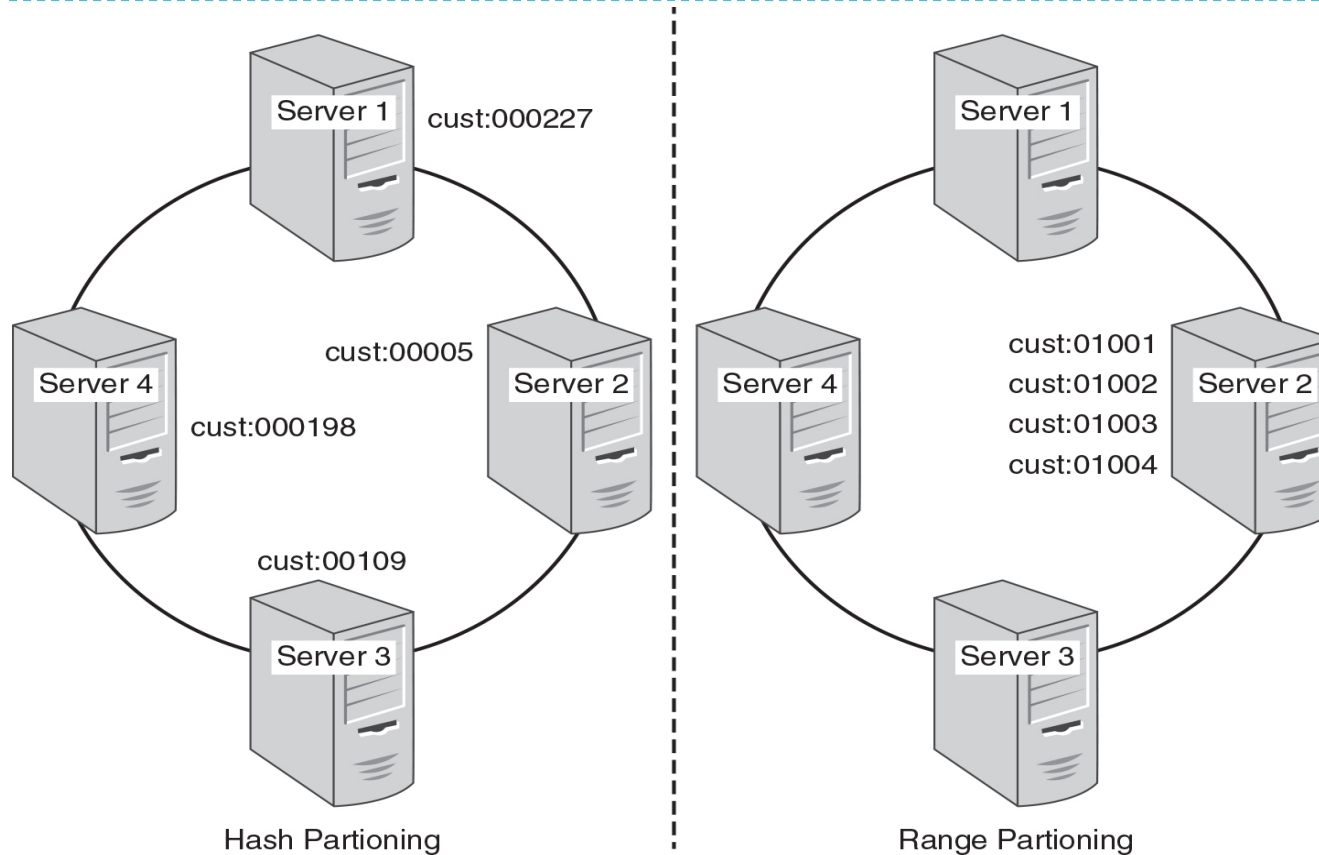
<pre>{   "PersonID": 101,   "LastName": "Smith",   "FirstName": "Fred",   "Phone": "555-4321" }</pre>
<pre>{   "PersonID": 102,   "LastName": "Jones",   "FirstName": "Mary",   "Address": {     "Street": "123 Main",     "City": "Anytown",     "State": "OH",     "ZIPCode": 12345   } }</pre>
<pre>{   "PersonID": 103,   "LastName": "Stephens",   "FirstName": "Howard",   "Address": {     "Street": "123 Main",     "City": "London",     "PostalCode": "ER3 5K8"   },   "FavoriteColor": "Blue" }</pre>

## Items

## attributes


<pre>{   "Artist": "No One You Know",   "SongTitle": "My Dog Spot",   "AlbumTitle": "Hey Now",   "Price": 1.98,   "Genre": "Country",   "CriticRating": 8.4 }</pre>
<pre>{   "Artist": "No One You Know",   "SongTitle": "Somewhere Down The Road",   "AlbumTitle": "Somewhat Famous",   "Genre": "Country",   "CriticRating": 8.4,   "Year": 1984 }</pre>
<pre>{   "Artist": "The Acme Band",   "SongTitle": "Still in Love",   "AlbumTitle": "The Buck Starts Here",   "Price": 2.47,   "Genre": "Rock",   "PromotionInfo": {     "RadioStationsPlaying": [       "KHCR",       "KQBX",       "WTNR",       "WJJH"     ]   },   "TourDates": {     "Seattle": "20150625",     "Cleveland": "20150630"   },   "Rotation": "Heavy" }</pre>
<pre>{   "Artist": "The Acme Band",   "SongTitle": "Look Out, World",   "AlbumTitle": "The Buck Starts Here",   "Price": 0.99,   "Genre": "Rock" }</pre>

# Common Partitioning Strategies Supported



- ▶ Note: Riak and DynamoDB offer Hash Sharding (Partitioning).
- ▶ DynamoDB - also offer a combination where a composite key is used
- ▶ i.e. Partition Key: Sort Key

# Mapping the Hash Key Value to a node



Could use **Modulo Arithmetic** to Determine the node the key-value will reside

Hash key is passed through a Modulo Function


If we have a hash key of 27 where will the Key Value be placed on the ring?

**BUT** - this works well UNTIL you **add** or **remove** nodes from the ring and become unworkable or serious amount of remapping of key-values to nodes is required.


**Consistent Hashing is the answer!**

Figure: An eight-server cluster in a ring configuration with modulo number assigned.

# Consistent Hashing: The Concept




- ▶ Node A, B, C and D in ring each with a token
- ▶ 5 Key Value objects say 1,2,3,4,5
- ▶ Each position in the ring represents hashCode value.



- ▶ **2 Scenarios –Removing and Adding nodes**
- ▶ Node C is removed
  - ▶ Object 5 now belongs on node D
- ▶ Node E is added
  - ▶ Object 4 now belongs on the node E
- ▶ The key value objects 4 and 5 are moved to the nodes

# Riak Key Value Store (Amazon DynamoDB inspired)

- ▶ Master-less – all nodes are equal
- ▶ Nodes can fail -> Hinted Hand-off
- ▶ Tunable Consistency offered
- ▶ Consistent hashing
  - ▶ Evenly distributes data across the cluster
  - ▶ No Manual sharding required
  - ▶ **160-bit integer keyspace**
  - ▶ Divided into fixed number of evenly sized partitions called **vnodes** as there is a vnode process that manages a partition
  - ▶ Partitions are claimed by nodes in the cluster
  - ▶ Replicas go to the N (=3 default) partitions following the key



Deployed as a cluster of physical nodes





Figure 1: The Riak "Ring"

# Hinted Handoff

- ▶ Allows Riak nodes to temporarily take over storage operations for a failed node and updates that node with changes when it comes back online.
- ▶ Used in other NoSQL databases like Cassandra
- ▶ Ensures High Availability





# Key-Value Stores - Key Features (1)

---

- ▶ Consistency
  - ▶ In distributed implementations using replication, **eventual consistency** supported
  - ▶ How could Write Conflicts (Entropy) be resolved?
    - ▶ Last Write wins OR
    - ▶ All values returned to client for conflict resolution ( manual resolution)
    - ▶ Riak uses Read Repair and later versions (1.3 onwards) use Active Anti Entropy ( via Merkle trees)
- ▶ No ACID Transaction support
  - ▶ Generally Speaking - No Guarantees on the writes particularly in distributed replicated environment – RIAK uses quorums for strong consistency
  - ▶ Partial Updates are not possible. An update is either a Delete or an Insert
  - ▶ Deletes – Using hashing –
    - ▶ Deleted values are marked as being deleted
    - ▶ Avoids losing existing values that previously collided with the deleted item's key
- ▶ Scaling
  - ▶ Horizontal Scaling using Replication and Sharding generally supported

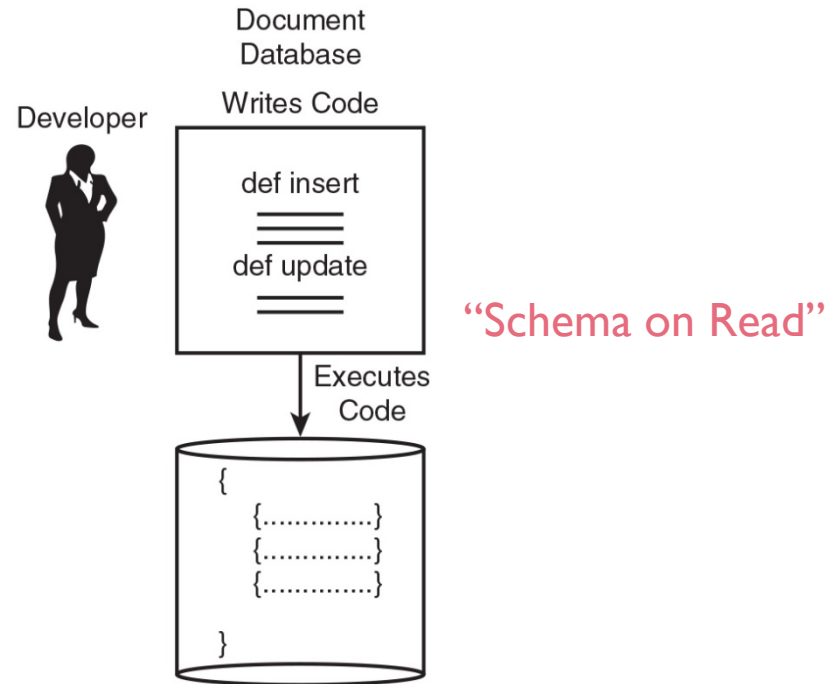
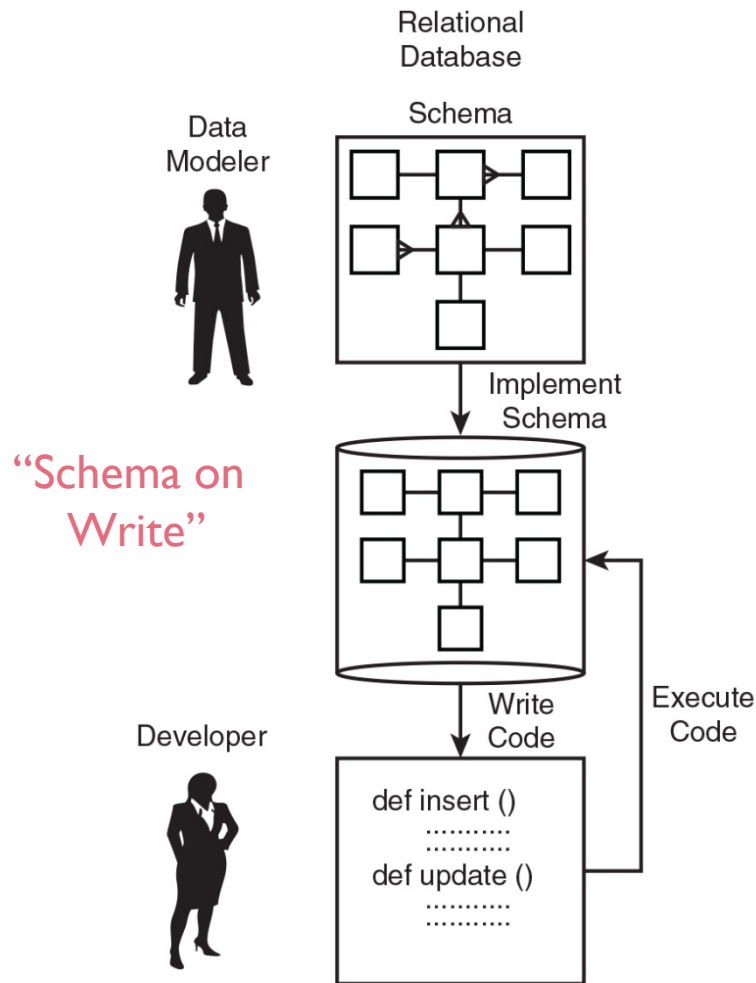
# Document Database

- ▶ A document database is similar in concept to a key/value store except that the values stored are documents.
- ▶ A document is a collection of named fields and values, each of which could be simple scalar items or compound elements such as lists and child documents.
  - ▶ Fields in a document can be encoded in a variety of ways, including XML, YAML, JSON, BSON, or even stored as plain text.
  - ▶ The fields in the documents are exposed to the database management system,
  - ▶ An application can query documents by using the document key (but a more common approach is to retrieve documents based on the value of one or more fields in a document)
  - ▶ Some document databases support indexing to facilitate fast lookup of documents based on an indexed field.
  - ▶ Many document databases support in-place updates
  - ▶ Usually Atomic at the Document Level
    - ▶ Read and write operations over multiple fields in a single document are usually atomic.





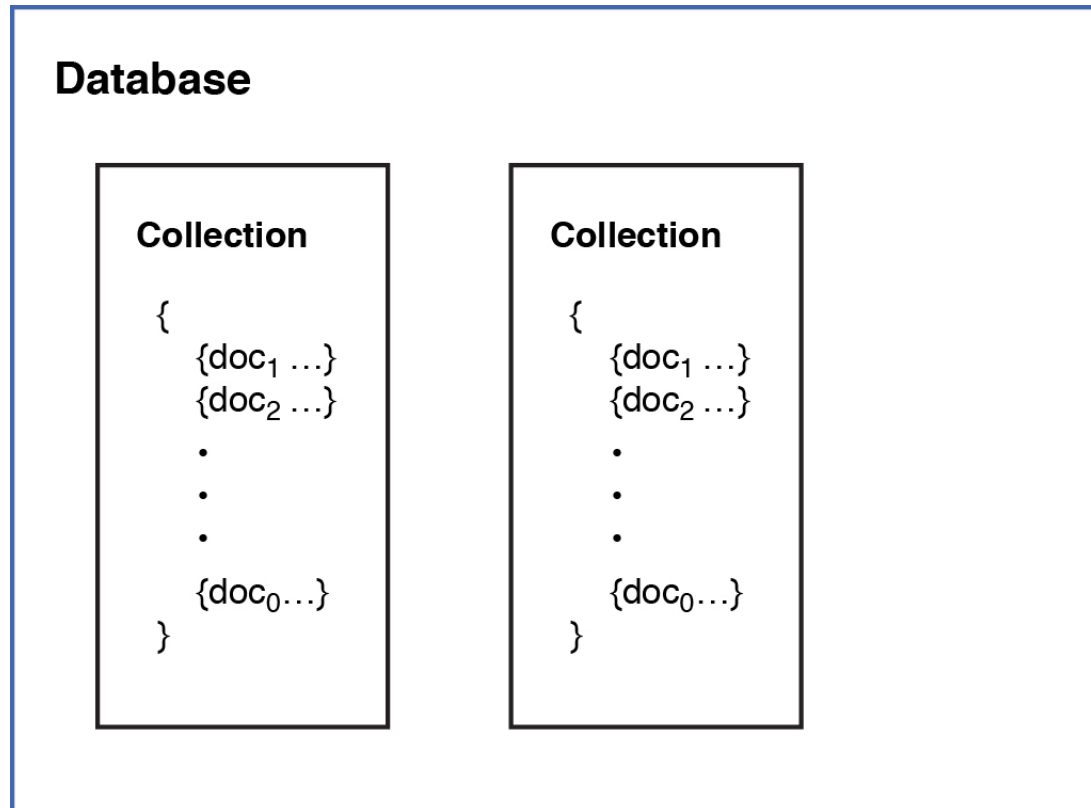
# NoSQL - Schemaless



Relational databases require an explicitly defined schema, but document databases do not. ( also applies to KV stores and Column Family stores)

# Document Database - Conceptual

---



The database is the highest-level logical structure in a document database and contains collections of documents.

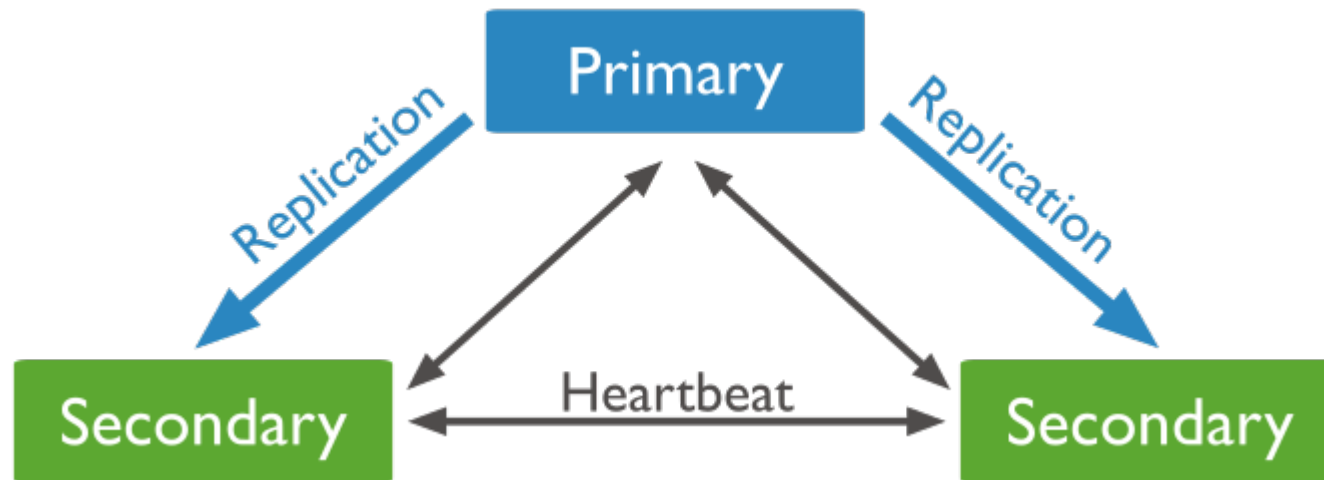


# MongoDB Terminology

Oracle	MongoDB
Database Instance	MongoDB Instance
Schema	Database
Table	Collection
Row	Document
Rowid	<u>_id</u>
Join ( or more correctly Foreign key)	No Joins- 2 <sup>nd</sup> Query; \$Lookup

# MongoDB Replica Sets

---



- ▶ A replica set is a group of [mongod](#) instances that maintain the same data set. A replica set contains several data bearing nodes and optionally **one arbiter node**. Of the data bearing nodes, one and only one member is deemed the **primary node**, while the other nodes are deemed **secondary nodes**.
- ▶ Source: MongoDB Reference Manual



# Key Features MongoDB(1)

---

- ▶ **Consistency**
  - ▶ Applicable for operations on a single document default
  - ▶ Uses Master/Slave replication
  - ▶ Eventual Consistency to Strong Consistency.
- ▶ **ACID Transactions not supported generally**
  - ▶ Transactions at the single-document level are known as atomic transactions (default)
  - ▶ Atomic Transactions across multiple documents now possible with version 4.04
  - ▶ **Durability and Consistent Writes**
    - ▶ Main levels of **WriteConcern** are offered
      - ❑ **0** => write operation not acknowledged
      - ❑ **1** => acknowledgement from the primary only required - DEFAULT
      - ❑ **Majority** => acknowledgement from majority in replica set required
      - ❑ **n** => acknowledgement from n members of replica set required
    - ▶ Main levels of ReadConcern offered
      - ❑ **“local”** – return the instances most recent data available even if not persisted to a majority of replica set members and may be aborted
      - ❑ **“majority”** – read data that has been written to a majority of replica set members and thus cannot be aborted( only with WiredTiger storage engine)



# Column-Family Database

---

- ▶ A column-family database organizes its data into rows and columns;
  - ▶ Aka *Wide Column databases* or *Extensible Record databases*
  - ▶ In its simplest form it can appear very similar to a relational database, at least conceptually.
  - ▶ The real power is its denormalized approach to structuring sparse data.
  - ▶ You can think of a column-family database as holding tabular data comprising of rows and columns, but you can divide the columns into groups known as *column-families*.
  - ▶ Each column-family holds a set of columns that are logically related together.
  - ▶ Can be schema-less **but can also define a schema!**

A P A C H E  
**HBASE**

Google™  
**BigTable**

 **Cassandra**



# Column Family Store: Cassandra

Oracle	Cassandra (peer-peer replication)
Database Instance	Cluster
Database Schema	Keyspace
Table	Column Family (aka Table)
Row	Row
Rowid	Row Key
Column ( same for all rows)	Column (can be different per row)



# Conceptual Structure of Data in a Column Family

CustomerID	Identity Column Family	
1	Title	Mr
	FirstName	Mark
	MiddleName1	William
	LastName	Hanson
2	Title	Ms
	FirstName	Lisa
	MiddleName1	Sarah
	MiddleName2	Louise
	LastName	Andrews
3	FirstName	Walter
	LastName	Harp

**A simple Column Family** with a set of column names

Row Key	Column Families	
CustomerID	CustomerInfo	AddressInfo
1	CustomerInfo:Title Mr CustomerInfo:FirstName Mark CustomerInfo:LastName Hanson	AddressInfo:StreetAddress 999 Thames St AddressInfo:City Reading AddressInfo:County Berkshire AddressInfo:PostCode RG99 922
2	CustomerInfo:Title Ms CustomerInfo:FirstName Lisa CustomerInfo:LastName Andrews	AddressInfo:StreetAddress 888 W. Front St AddressInfo:City Boise AddressInfo:State ID AddressInfo:ZipCode 54321
3	CustomerInfo:Title Mr CustomerInfo:FirstName Walter CustomerInfo:LastName Harp	AddressInfo:StreetAddress 999 500th Ave AddressInfo:City Bellevue AddressInfo:State WA AddressInfo:ZipCode 12345

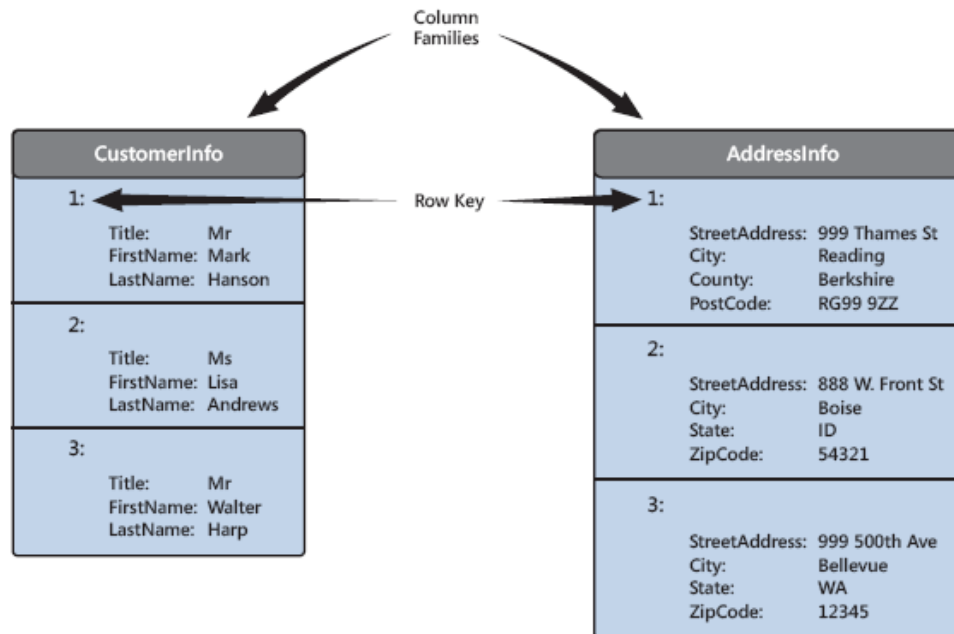
Different rows can have different fields and the data does not conform to a rigid layout.

From a physical perspective, Column Families are generally stored separately



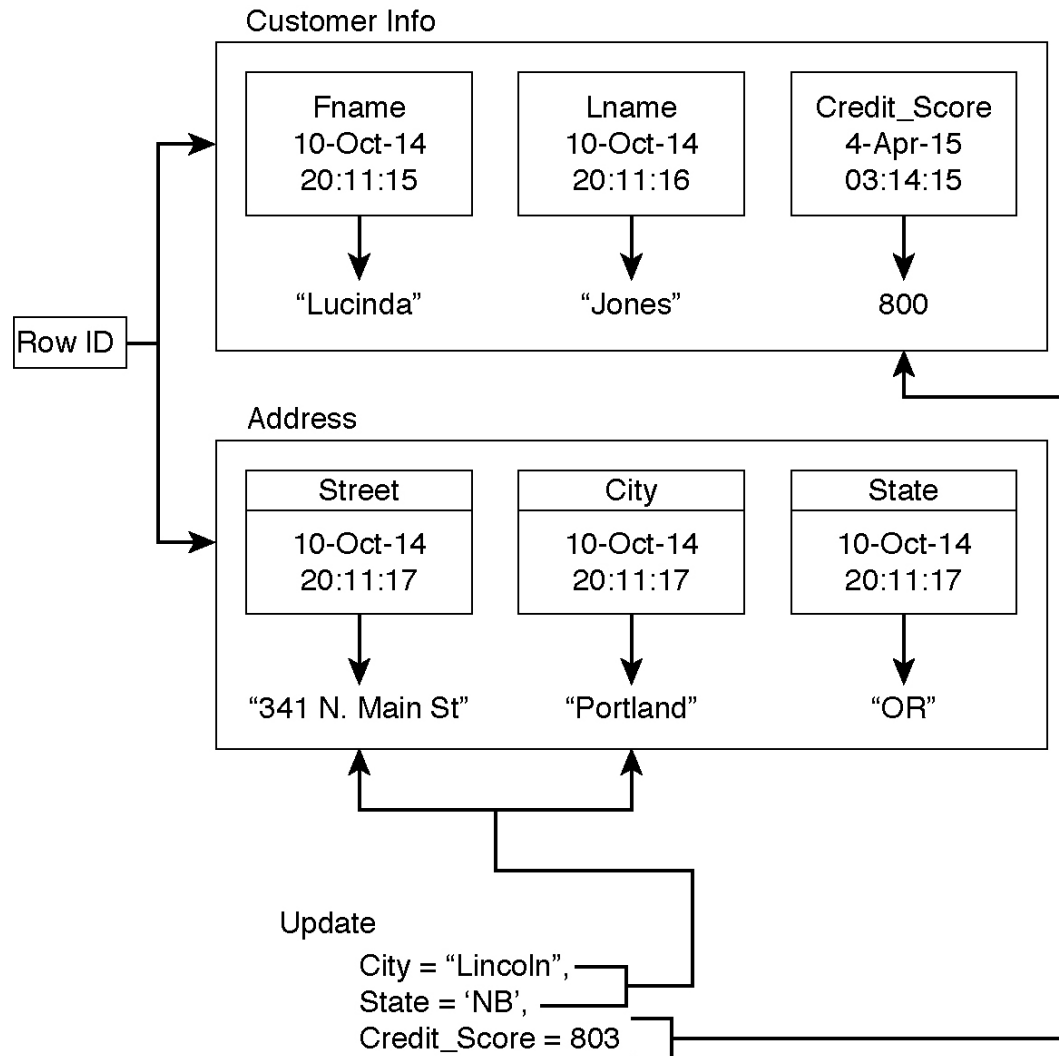
# Structure of Data in a Column Family contd

- ▶ The data for a single entity that spans multiple column-families will have the same row key in each column-family.



- ▶ In a single column-family, the data for the part of a row in a single column-family are usually atomic, although a number of implementations provide atomicity across the entire row (spanning multiple column-families e.g. Hbase, BigTable) as well.

# Atomicity: Google BigTable An Example



Read and write operations **are atomic at the Row Level**. All columns are read or written or none are.

Note: values are not Overwritten but appended

# Document DB VS Column Family DB

## Document Database:

{fname: 'Lucinda',  
lname: 'Jones',  
Credit\_Rating: 800,

Address:

{Street: '341 N. Main St.',  
City: 'Portland',  
State: 'OR'  
}

{fname: 'Frank',  
lname: 'Antonio',  
Credit\_Rating: 768  
}

## Column Family Database:

Customer\_Info

fname	lname	Credit_Rating
10-Oct-14 20:11:15	10-Oct-14 20:11:16	04-Apr-15 03:14:15
↓	↓	↓
"Lucinda"	"Jones"	800
fname	lname	Credit_Rating
24-May-13 07:01:01	24-May-13 07:01:01	24-May-13 07:01:02
'Frank'	'Antonio'	768

Row ID<sub>1</sub>

Row ID<sub>2</sub>

Address

Street	City	State
10-Oct-14 20:11:15	10-Oct-14 20:11:16	10-Oct-14 20:11:17
"341 N. Main St."	"Portland"	'OR'



# Wide Column Family Example

---

StockTicker	StockPrices Column Family	
ABBT	01/01/2013 11:54:16	130
	01/01/2013 11:58:22	131.5
	01/01/2013 12:02:58	132
	01/01/2013 12:03:18	135
	01/01/2013 12:08:57	133.5
	...	
BAXD	01/01/2013 11:58:32	11.5
	01/01/2013 11:59:42	10.5
	01/01/2013 12:08:30	9
	01/01/2013 12:09:24	9.5
	...	
EFCD	01/01/2013 12:01:17	55
	01/01/2013 12:07:12	54.5
	01/01/2013 12:10:32	55
	01/01/2013 12:19:14	55.5
	01/01/2013 12:20:40	56
	...	
FAMI	01/01/2013 11:57:19	228
	01/01/2013 12:09:45	225
	01/01/2013 12:15:48	227
	01/01/2013 12:19:55	27.5
	01/01/2013 12:22:09	228.5
	...	

What are the  
column names?

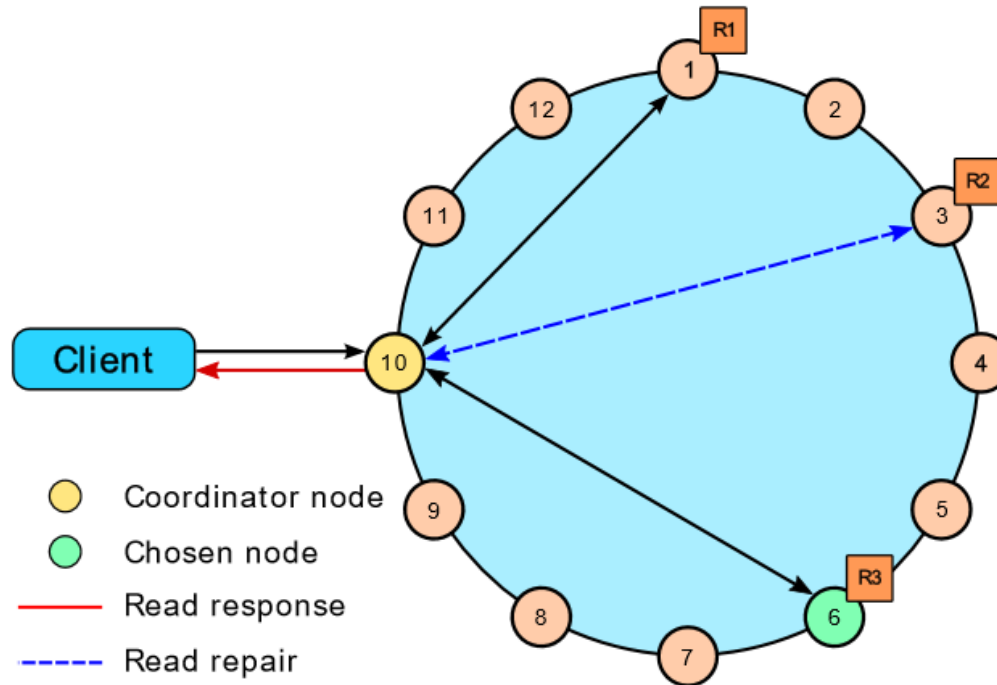
What is the Row  
Key?

Example of a Dynamic Column Family

Column Families can contain  
thousands of columns.

# Cassandra

## Logical Ring – **Read Request** – an example



3 Replica Nodes with a consistency level set to QUORUM (majority = 2 here)


Returns the value with most recent timestamp

- Logical Ring of vNodes
- Gossip protocol use for inter node communications
- Tunable Consistency supported
- Supports **Consistent Hashing** as well as **Read Repair**
- Supports **Hash Partitioning** (default) and Range Partitioning
- CQL, a SQL like language has become the primary API


# Cassandra

## Logical Ring – **Write Request An Example**


- ▶ **Write Consistency Level of ONE**
  - ▶ Replicas will eventually be updated
  - ▶ There are other Consistency Levels
- ▶ Ensure that the value is written to one replica to be deemed successful




# Gossip Protocol




2 Messages



6 Messages



12 Messages



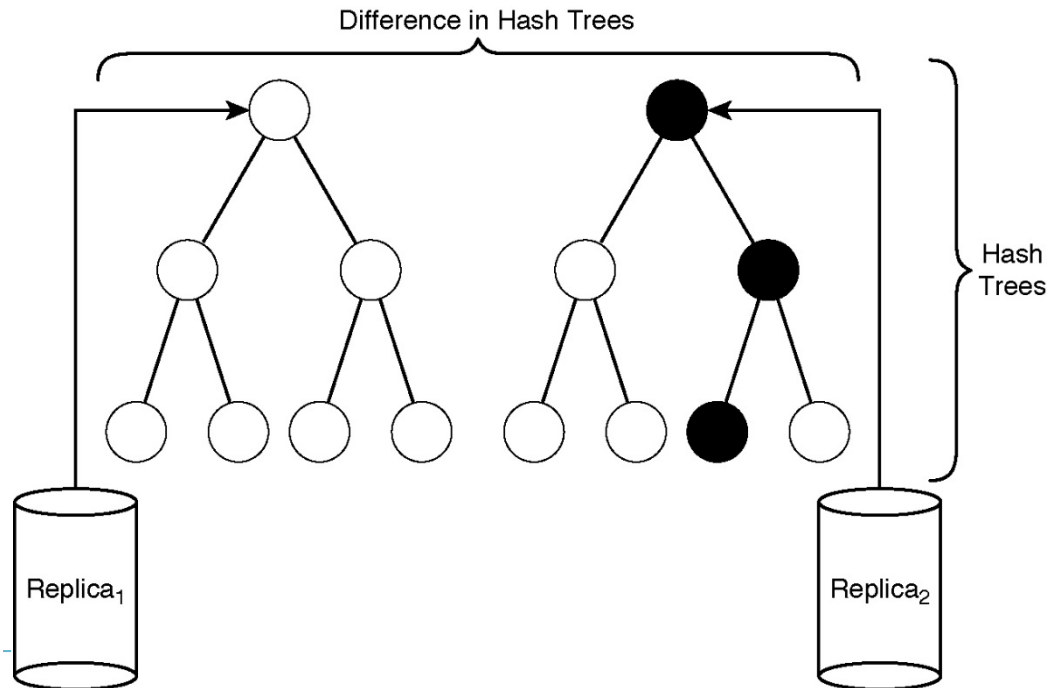
20 Messages

The number of messages sent in a complete server-to-server communication protocol grows more rapidly each time a server is added to the cluster.

- ▶ If  $N$  is the number of servers the  $N \times (N-1)$  is the number of messages needed to update all servers with information about all other servers! Very Expensive!
- ▶ A more efficient method of sharing information is the Gossip Protocol

# Anti-Entropy

- ▶ The process of detecting differences in replicas
- ▶ 2 common Methods
  - ▶ Read Repair
  - ▶ Hash trees, or Merkle trees,
    - ▶ Allow for rapid checks on consistency between two data sets.





# Key Features

---

## ▶ Querying

- ▶ Can query particular column family rather than all column families that make up a row
  - ▶ Remember – Key/Value Stores often returns the whole entity

## ▶ Transactions

- ▶ Piece-wise updates supported
- ▶ Not the best solution for systems that require full ACID transactions
- ▶ Atomicity at the Row level
  - ▶ Cassandra – atomic at the row level i.e. a single read/write will either succeed or fail

## ▶ Cassandra Peer-Peer Logical Ring for read and writes

## ▶ Sharding\Partitioning and Replication supported

- ▶ Eventual Consistency as well as Stronger Consistency supported where all replicas must respond for transactions to be successful.
  - Read and Write Quorums also supported

# Key Features

---

## ▶ Appropriate when

- ▶ Real-time random\read access capability is required and data being stored has some defined structure
- ▶ Applications that that require the ability to **always write**
- ▶ Applications with dynamic fields and can tolerate short-term inconsistencies in replicas
- ▶ Applications with **truly large volumes** of data (**100's ofTB**)
- ▶ Applications that are geographically distributed over multiple data centres

## ▶ Not Appropriate when

- ▶ Joins are required
- ▶ Systems that require ACID transactions
- ▶ Binary data needs to be supported

# Graph Databases

Graphs have a history dating back to 1736, when Leonhard Euler solved the “Seven Bridges of Königsberg” problem

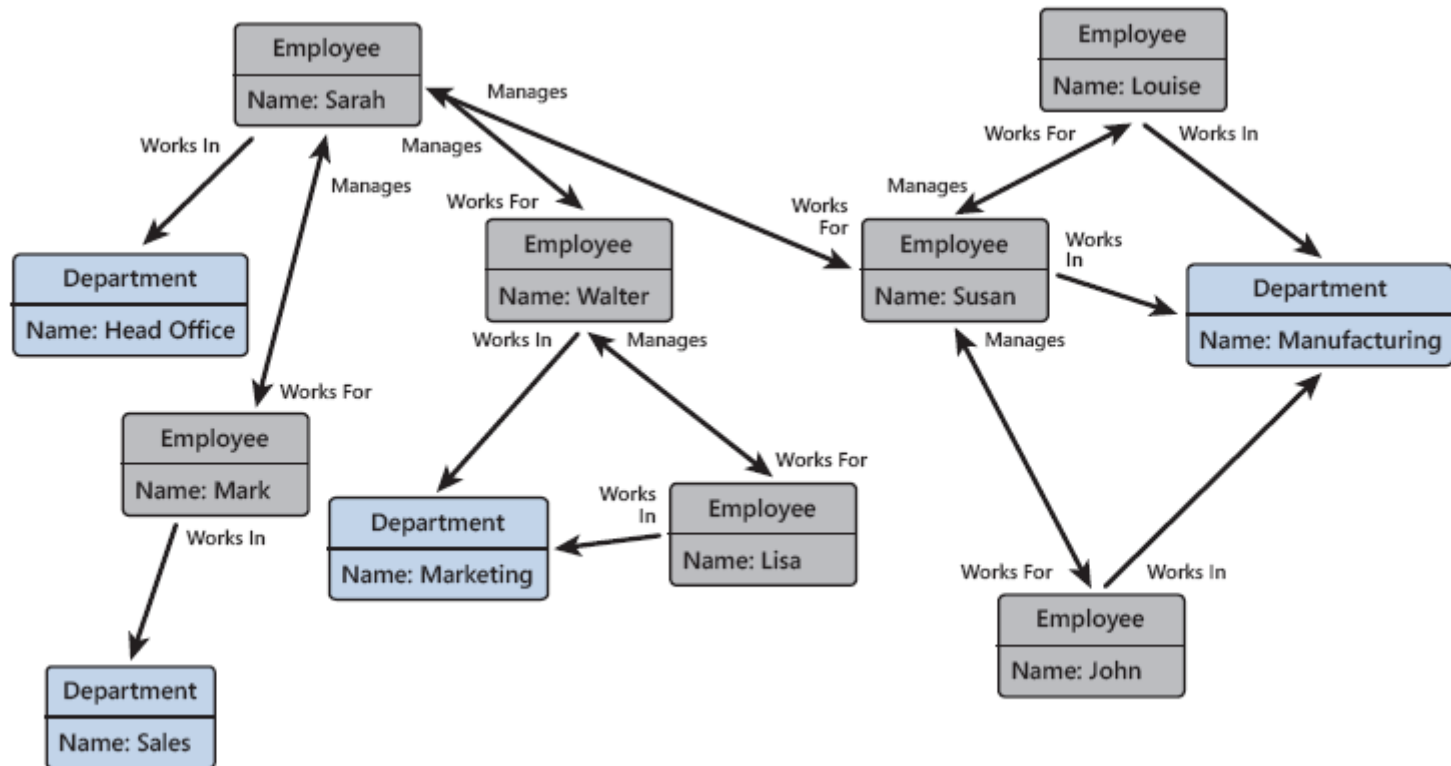
- ▶ Stores two types of information:
  - ▶ **Nodes** (aka vertices) that you can think of as **instances of entities**, and
  - ▶ **Relationships** (aka links or edges) which specify the relationships between nodes.
  - ▶ Nodes and edges can both have properties that provide information about that node or edge (like columns in a table). Additionally, edges can have a direction indicating the nature of the relationship.
- ▶ The purpose of a graph database is to enable an application to efficiently perform queries that traverse the network of nodes and edges, and to analyze the relationships between entities.



**FlockDB**



# HR Data Structured as a Graph





# Comparing with Relational Database

---

## ▶ Pros

- ▶ Schema Flexibility
- ▶ More intuitive querying
- ▶ Avoids Joins
- ▶ Local Hops are not a function of total nodes
- ▶ ACID Compliant

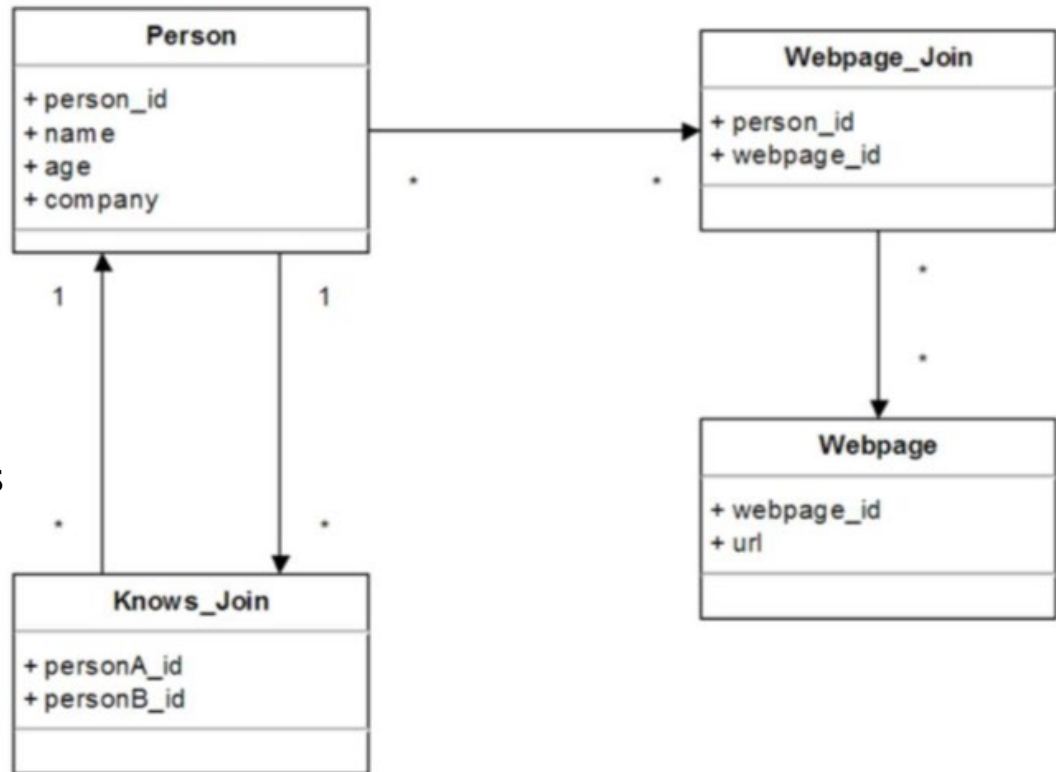
## ▶ Cons

- ▶ Query Languages are not unified



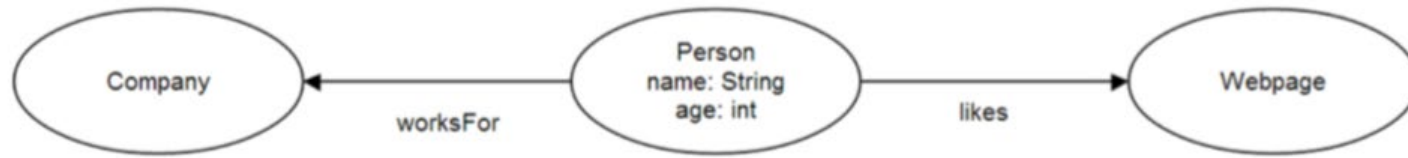
# Relational Database Example

- ▶ In this example, we have a schema of a relational database that stores people, their name, age, the company they work at, who they know, and webpages that they like.
- ▶ The **knows** relation and Webpage(**likes**) relation is represented by join tables.
- ▶ Perhaps due to poor oversight, this schema only stores one company per person.
- ▶ What if we wanted to allow multiple companies? We could either allow duplicates rows (!) or create another join table.




# Graph Database Example

---




- ▶ This diagram is the graph database schema that represents the same people data that was presented before. The circles represent nodes (vertices), and the solid lines represent relationships (edges).
- ▶ To solve the same problem in a graph database, we need only create a new edge from the Person node to the Company node. This is a much simpler solution.
- ▶ An important thing to note is that graph databases do not have to execute joins for each edge traversal, and avoid join bombs!

# Neo4j – offers a form of Master Slave





# Cache Sharding supported in Neo4j





# Features(1)

---

## ▶ Consistency

- ▶ Graph Database usually do not support distributing nodes on different servers
  - ▶ Within single server data is always consistent
  - ▶ Neo4j ( open source, java) – fully ACID Compliant on a single server environment

## ▶ Transactions

- ▶ Transactions are supported
  - Neo4j must initiate a transaction for a Write
  - fully ACID Compliant on a single server environment

## ▶ Availability

- ▶ Neo4j provides high availability through replicated slaves
- ▶ Slaves can also handle writes but the write must be committed at the master and then at the slave

## ▶ Some Scaling options available

## ▶ Querying

- ▶ Indexing supported
- ▶ A number of query languages are supported
  - Gremlin; Domain specific language for traversing graphs; Neo4J - Cypher

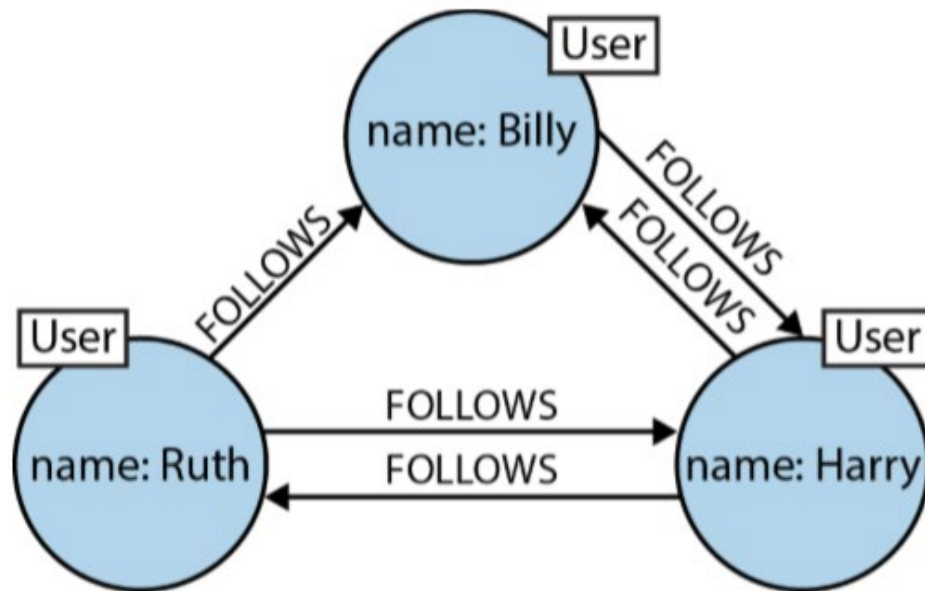
# FOUR building Blocks in Neo4j Graph Database

---


- ▶ Nodes
- ▶ Relationships
- ▶ Properties
- ▶ Labels




## Nodes & Relationships



# Relationships



**Nodes can have more than one relationship**



**Nodes can be connected by more than one relationship**



**Self relationships are allowed**

# Relationships


---

- ▶ Every relationship has a name and a direction
- ▶ Can contain properties
- ▶ Used to represent quality or weight of relationship or metadata
- ▶ Every relationship must have a startnode and endnode
- ▶ Relationships are defined with regard to node instances ,not classes of nodes
- ▶ Relationship=>A join in a relational database
- ▶ Two nodes representing the same kind of “thing” can be connected in very different ways

# Labels

---

- ▶ Every node can have zero or more labels
  - ▶ Node instance Ruth has a label User but could also have another label Author for example
- ▶ Used to represent roles (e.g.user, product, company)
- ▶ Group nodes
- ▶ Allow us to associate indexes and constraints with groups of nodes



# Cypher Query Language

---

- ▶ Declarative Pattern-Matching language
- ▶ SQL-like syntax
- ▶ Designed for graphs
- ▶ E.g.




**MATCH (a)-->(b) RETURN a, b**

--Returns all nodes and their relationships in the graph database (a and b are just placeholders)



# Paths -Examples

---



-- Looks for this pattern

```
MATCH (a)-[:ACTED_IN]->(m)<-[:DIRECTED]-(d)
```

```
RETURN a.name, m.title, d.name;
```

Provide actor's name, movie title and director' name to output returned

# Paths - Another way

---




```
MATCH (a)-[:ACTED_IN]->(m), (d)-[:DIRECTED]->(m) RETURN a.name, m.title, d.name;
```

An Implied AND



Paths – returns the nodes that are related by a **ACTED IN** and **DIRECTED** relationships

---



```
MATCH p=(a)-[:ACTED_IN]->(m)<-[:DIRECTED]-(d)
RETURN nodes(p);
```



# Node Syntax

---

- ▶ **()**
  - ▶ #represents an anonymous, uncharacterized node
- ▶ **(matrix)**
  - ▶ #if we want to refer a node we need an identifier
- ▶ **(:Movie)**
  - ▶ #represents anonymous, characterized node
- ▶ **(matrix:Movie)**
  - ▶ #represents a referenced, characterized node i.e. matched Movie nodes
- ▶ **(matrix:Movie {title: "The Matrix"})**
  - ▶ #Movie node with a particular property value
- ▶ **(matrix:Movie {title: "The Matrix", released: 1999})**
  - ▶ #Movie node with a number of property values



# Relationship Syntax

---

- ▶ Cypher uses a pair of dashes (--) to represent an undirected relationship.
- ▶ Directed relationships have an arrowhead at one end (eg, <--, -->).
- ▶ Bracketed expressions (eg: [...]) can be used to add details.
- ▶ This may include identifiers, properties, and/or type information.
- ▶ Similar Concepts to the Node Syntax

--> #represents an anonymous, uncharacterized node

-[role]-> #if we want to refer a relationship we need an identifier\alias

-[:ACTED\_IN]-> #represents anonymous characterized relationship

-[role:ACTED\_IN]-> #represents a referenced, characterized relationship

-[role:ACTED\_IN {roles: ["Neo"]}]> #ACTED\_IN relationship with a particular property value